
Deep Learning Based Superconductivity Prediction and Experimental Tests

Daniel Kaplan¹ Adam Zheng² Joanna Blawat³ Rongying Jin³

Viktor Oudovenko¹ Gabriel Kotliar¹ Weiwei Xie² Anirvan M. Sengupta^{1,4}

¹ Department of Physics and Astronomy, Rutgers University, Piscataway, NJ 08854

² Department of Chemistry, Michigan State University, East Lansing, MI 48824

³ Department of Physics and Astronomy, University of South Carolina, Columbia, SC 29208

⁴ CCM and CCQ, Flatiron Institute, New York, NY 10010

d.kaplan1@rutgers.edu, dzheng_box@hotmail.com, jblawat@email.sc.edu,
rjin@mailbox.sc.edu, udo@physics.rutgers.edu, kotliar@physics.rutgers.edu,
xieweiwe@msu.edu, anirvans.physics@gmail.com

Abstract

The discovery of novel superconducting materials is a longstanding challenge in materials science, with a wealth of potential for applications in energy, transportation, and computing. Recent advances in artificial intelligence (AI) have enabled expediting the search for new materials by efficiently utilizing vast materials databases. In this study, we developed an approach based on deep learning (DL) to predict new superconducting materials. We have synthesized a compound derived from our DL network and confirmed its superconducting properties in agreement with our prediction. Our approach is also compared to previous work based on random forests (RFs). In particular, RFs require knowledge of the chemical properties of the compound, while our neural net inputs depend solely on the chemical composition. We further discuss the existing limitations and challenges associated with using AI to predict and discover new superconductors, along with potential future research directions.

1 Introduction

Superconductivity is a phenomenon characterized by zero electrical resistivity and the Meissner effect – the diamagnetic expulsion of magnetic fields from the bulk of the sample, and is microscopically attributed to the formation of Cooper pairs, making it an ideal system for the study of quantum entanglement [2, 14, 3]. The search for new superconducting materials has been a long-standing challenge in materials science with potential applications in energy, transportation, and computing [26, 35, 12]. Traditional methods of discovering new superconductors involve extensive experimental trial and error, which can be time-consuming and expensive. In recent years, researchers have explored the use of machine learning algorithms to accelerate the discovery of new materials [1, 21, 16, 24].

Modern machine learning techniques, such as deep learning [19], combined with genetic algorithms [10], can efficiently screen large databases of materials properties and predict the properties of potential new materials [30, 29, 34]. Beyond identifying potential materials, machine learning has also emerged as a tool for providing researchers with a deeper understanding of certain properties. For example, machine learning was used to develop simple models that explain the electronic structures of half-Heusler phases for potential thermoelectric properties [11], or direct the search for compounds with high hardness [23] by predicting the elastic moduli as a proxy [22], and so on.

In this work, we evaluated the effectiveness of deep learning methods to classify materials into superconductor/non-superconductors and also to predict T_c when it is a superconductor. Our tools were applied to predict superconducting properties of hypothetical materials in sigma phases and compared to experimental findings. On this very special and small set of materials, the performance of both the random forest method and our deep learning networks are mixed. Still, we believe that we are at the beginning of the integration of AI and experimental synthesis, and this integration has immense potential for expediting the discovery of new superconductors and facilitating a deeper understanding of the underlying superconductivity phenomena. Previous work has employed regression-type model [33] or, in the realm of DNN (deep neural networks), neural networks for T_c prediction without classification [18], while Deep Set architecture was used was classification in Ref. [28]. Recently, superconductor design through gradient optimization [13] has been used. Gradient boosting has also been employed, e.g. in Refs. [15, 20]. Our approach differs from these, as we present below. We also find superior performance to regression algorithms trained with chemical features. In the present work, we rely solely on the chemical formula of the compound.

2 Deep Learning Superconductivity Prediction

2.1 Data and Preprocessing

The superconducting data for alloys were retrieved from database for superconducting materials, SuperCon [25]. The data tabulates the experimental results of measurements of superconductivity for 16,414 compounds. Each row indicates the elements in the compound and their respective percentages (the input), as well as the measured T_c value (the output), where $T_c > 0$ indicates that the compound is superconducting.

For each data point, a data vector of length 120 is generated with the index representing the atomic number of elements and its corresponding percentage being the entry for that index. These generate 16,414 sparse 120-dimensional vectors. This will be the input of our fully connected network. In principle, it may be useful to incorporate some information about the relative position of elements in the periodic table. Several works [18, 13] break up the periodic table into multiple rectangles and input the chemical compound as multiple images which is then processed by a convolutional network (CNN). We will reshape the 120 dimensional vector into a single 10×12 ‘image’. This approach is harder to justify, a priori, but, as we will see, our CNN seems to have slightly better performance in some cases, compared to the fully-connected network. We speculate that it benefits from bringing together chemically similar elements.

The output comprises two components: classification ground truth, represented by a single binary value of 0 or 1, indicating the presence or absence of superconductivity. Thus, the two classes correspond to non-superconducting compounds (0) and superconducting compounds (1). The other component is the measured T_c value.

2.2 Architecture

Given the nature of the dataset, we decided to treat the problem as a combination of classification and regression tasks. In other words, we construct deep neural nets (DNNs) that can classify whether a given compound is superconductive or not and predict the corresponding T_c value. The DNN model comprises three components: the network backbone and two prediction branches: one for classification and the other for T_c value. The architecture is schematically illustrated in Fig. 1, left. We build two networks, one fully-connected, the other convolutional. For the details of the architectures, please see Supplementary Materials. The DNN model is implemented using the PyTorch [27] framework, which accepts network data in the 4-D tensor format of NCDHW (batch N, channels C, depth D, height H, width W).

A word of justification for our choice of the branched architecture and on our planned method of training. Since the two tasks, namely, the classification of superconducting nature and the regression of T_c are related, it makes sense to train a set of shared latent variables as neurons in the initial layers. Since the teaching signal from T_c is more specific than the signal from the class, we train the shared part via the regression task. Since the classification task is conceptionally easier than the regression task, we choose to allow some additional flexibility to the network, rather than restricting

to it classifiers obtained by thresholding the output of a regression network, as in Refs. [33, 18] and others.




Figure 1: Left: DNN structure with a shared backbone and two prediction branches, one for the T_c value and the other for classification. We have used our DNN to predict a new superconductor (see Sec. 3) from the σ family of alloys, and measured its resistivity as a function of temperature, extracting T_c .

2.3 Training

The data set is divided into 2 portions: training (13132 entries) and testing dataset (3282 entries). During the training process, the Mean Squared Error (MSE) is utilized to compute the loss error, which is subsequently leveraged for the backpropagation calculation using the Adam optimizer [17]. The training process is carried out in two stages. First, the Backbone and T_c prediction branches are jointly trained, and the loss of the predicted T_c value versus the ground truth is used for network backpropagation. Once the T_c Prediction results reach a satisfactory level, the parameters of the Backbone and T_c Prediction branches are fixed. The training is then performed solely on the Classification branch. Each branch was run for 5,000 epochs. This procedure was performed for both architectures. The learning curves, showcasing performance in accuracy and MSE loss are plotted in Fig. 2

After completing Stage1 and Stage2 training, the model is finalized for testing. The training process is further optimized by using two-step learning rate decay. The initial learning rate is set to 0.0001, and at 3000 epochs, it is decayed to 0.00001 to fine-tune the training parameters. The training results are evaluated based on the backbone and T_c prediction branch training, and the T_c prediction difference is calculated by averaging the mean absolute error (MAE) difference between the predicted and ground-truth values [8]. A lower difference value indicates more accurate prediction results. Both the loss and T_c prediction differences decrease with an increase in epoch count, and starting from




Figure 2: Performance against epochs of training for the CNN (convolutional neural network) and FC (fully-connected neural network). See main text for details.

the 2000th epoch, the loss and difference curve reach a plateau. However, at the 3000th epoch, with a decay in the learning rate from 0.0001 to 0.00001, the loss and T_c prediction differences further decrease, indicating that the learning rate decay indeed helps to improve the training process by providing finer granularity control of parameter updating.

2.4 Results

Results and the performance of the classification branch is evaluated based on its accuracy in correctly identifying if a compound exhibits superconductivity or not, as compared to the ground-truth. A higher accuracy value indicates better performance in the classification task. The testing process involves evaluating the performance of the trained model on a separate dataset known as the test dataset, which is distinct from the data used during model training. The test dataset comprises a total of 3282 entries. The final training and testing results with appropriate comparison to random forest methods are listed in Table 1.

Table 1: The summary of classification test results for various methods. RF= Random Forest [33], FCNN= Fully-connected neural net, CNN= Convolutional Neural Net. The error in T_c regression (Reg.) is defined as the mean of $T_c(\text{pred.}) - T_c(\text{actual})$.

Model	Accuracy	Precision	Recall	F1	Reg. [K]	Class.
RF [33]	85% ¹	85%	94%	90%	7.88 ± 0.52	NA ²
FCNN	$83 \pm 0.6 \%$	$85 \pm 0.7\%$	$95 \pm 0.4\%$	90%	4.497 ± 0.328	$83.04 \pm 0.6\%$
CNN	$85 \pm 0.3\%$	$88 \pm 0.4\%$	$92 \pm 0.3\%$	90%	3.208 ± 0.180	$84.77 \pm 0.3\%$

The SuperCon database [25] has 16,414 entries, 12,499 ($\approx 76\%$) of which are known to be superconductors and the rest are considered non-superconducting. Note that a positive/negative weight guessing random classifier could achieve a maximum of $\approx 76\%$ accuracy by always going with the majority class [4]. In this dataset, the majority class predictor would have 100% recall but $\approx 76\%$ precision. These baselines are to be compared with the classification performances of the machine learning methods.

3 Experimental Studies

Superconductivity was observed in a few materials belonging to the σ -phase category [6], characterized by a high degree of disorder and the presence of multiple elements occupying the same Wyckoff position. The occurrence of σ -phases is strongly dependent on the concentration and the type of the constituent elements. Notably, superconductivity has been demonstrated in materials such as Nb-Ru-Ge, Nb-Rh-Ge, and Nb-Rh-Si, which has stimulated further investigations into Re-based materials with varying ternary systems to explore their superconducting properties [5, 7].

We leveraged our chemical intuition and expertise, to narrow our search for new superconductors to ternary compounds with formula unit $\text{Mo}_{20}\text{X}_6\text{Z}_4$, with $\text{X}=\text{Re, Rh, Ru}$, $\text{Z}=\text{Ge, Si}$, as the 20-6-4 stoichiometry is most likely to form in the sigma phase. Then we use our DNN-based algorithm to investigate the potential superconducting properties of these compounds. Our CNN predicted that **$\text{Mo}_4\text{Re}_2\text{Si}$** is likely to exhibit superconductivity with a critical temperature (T_c) of approximately 6 K, while **$\text{Mo}_4\text{Re}_2\text{Ge}$** is also predicted to be a superconductor with $T_c \sim 5$ K. The drop in resistivity for the associated, fully stoichiometric compound of the 20 – 6 – 4 family, **$\text{Mo}_{20}\text{Re}_6\text{Si}_4$** , is shown in Fig. 1(right). Subsection A.4 provides further details of the experimental set up, synthesis, crystal structure and additional measurements performed to confirm the 3D nature of the superconductivity.

Data leakage related issues are discussed in subsection A.3. For clarity, by starting with the loading composition of 4-2-1 ratio, we get final products with around 20-6-4 ratio.

¹ To match with our definition of accuracy, this value was extrapolated from Stanev *et al.* to the point $T_{sep} = 0$.

² Stanev *et al.* do not explicitly train for classification, we do not present a direct comparison as it would be fine-tuned by their value of T_{sep} .

4 Conclusion

This work proposes a DNN-based superconductivity prediction model that uses 13132 data points to achieve an accuracy of 84% on 3282 test data. This is a small dataset compared to many cutting edge applications of DNNs. For instance, popular image classification algorithms such as mobile-net and res-net that use ImageNet as a basis for training are trained on 1.2 million images to classify 1000 objects. The top-1 accuracy for ImageNet has moved into 90% plus territory only within the last few years (see <https://paperswithcode.com/sota/image-classification-on-imagenet>). Our performance is comparable to previous work using random forests, but does not require detailed atomic-level chemical information utilized by Stanev *et al.* [33] while using random forests. Unlike Konno *et al.* [18] and Pereti *et al.* [28], our dataset contains only real compounds extracted from SuperCon, without padding with fictitious entries for increased sensitivity. Our data is further not restricted to the high- T_c regime as in Ref. [13], but is composed of data for a wide range of superconducting temperatures. In classification we find comparable performance relative to a more involved algorithm as in Ref. [28].

5 Limitations of This Study and Future Directions

We train our model on a relatively small dataset with an abundance of superconductors. To enhance the prediction accuracy and robustness of the model, more training data, with appropriate materials distribution, is required. More importantly, the input in our dataset only contains the chemical composition of the material. One extension is to encode the neighborhood of elements in the periodic table, as in the graph sometimes known as the “periodic spiral” [31]. The sigma phase experiments indicate that all the methods have difficulties differentiating between elements in the same group. Utilizing the crystal structure [32] and electronic structure can provide additional information to the model to abstract during the training process. Also interesting is the use of symmetries for prediction [36].

These augmentations of predictors could boost the model’s performance and make predictions in real-world scenarios.

6 Acknowledgements

DK acknowledges the support of the Abrahams Postdoctoral Fellowship of the Center for Materials Theory, Rutgers University, and the Zuckerman STEM Fellowship. GK acknowledges the support of NSF DMR-1733071. WX’s work is supported by U.S.DOE-BES under Contract DE-SC0023648.

7 Data Availability

The original SuperCon data used in this study are available at the references given in this paper [25]. Datasets generated during the current study are available from the corresponding author on reasonable request. The code used in this work is freely available at: <https://github.com/danielkaplan137/TcMLPred.git>

References

- [1] Patrick Avery, Xiaoyu Wang, Corey Oses, Eric Gossett, Davide M Proserpio, Cormac Toher, Stefano Curtarolo, and Eva Zurek. Predicting superhard materials via a machine learning informed evolutionary structure search. *npj Computational Materials*, 5(1):1–11, 2019.
- [2] John Bardeen, Leon N Cooper, and John Robert Schrieffer. Theory of superconductivity. *Physical review*, 108(5):1175, 1957.
- [3] Charles P Bean. Magnetization of high-field superconductors. *Reviews of modern physics*, 36(1):31, 1964.
- [4] Jason Brownlee. *Probability for machine learning: Discover how to harness uncertainty with Python*. Machine Learning Mastery, 2019.

- [5] Elizabeth M Carnicom, Weiwei Xie, Tomasz Klimczuk, and Robert J Cava. New σ -phases in the Nb–X–Ga and Nb–X–Al systems (X= Ru, Rh, Pd, Ir, Pt, and Au). *Dalton Transactions*, 46(41):14158–14163, 2017.
- [6] Elizabeth M Carnicom, Weiwei Xie, Zuzanna Sobczak, Tai Kong, Tomasz Klimczuk, and Robert J Cava. Superconductivity in the Nb–Ru–Ge σ phase. *Physical Review Materials*, 1(7):074802, 2017.
- [7] Elizabeth M Carnicom, Tai Kong, Tomasz Klimczuk, and RJ Cava. The σ -phase superconductors Nb₂₀. 4Rh₅. 7Ge₃. 9 and Nb₂₀. 4Rh₅. 7Si₃. 9. *Solid State Communications*, 284:96–101, 2018.
- [8] Tianfeng Chai and Roland R Draxler. Root mean square error (RMSE) or mean absolute error (MAE)?—arguments against avoiding rmse in the literature. *Geoscientific model development*, 7(3):1247–1250, 2014.
- [9] Alexis Cook. Data Leakage. <https://www.kaggle.com/code/alexisbcook/data-leakage>, 2023. [Online; accessed 9-July-2024].
- [10] Omid E David and Iddo Greental. Genetic algorithms for evolving deep neural networks. In *Proceedings of the companion publication of the 2014 annual conference on genetic and evolutionary computation*, pages 1451–1452, 2014.
- [11] Maxwell T Dylla, Alexander Dunn, Shashwat Anand, Anubhav Jain, and G Jeffrey Snyder. Machine learning chemical guidelines for engineering electronic structures in half-heusler thermoelectric materials. *Research*, 2020.
- [12] SR Foltyn, L Civale, JL MacManus-Driscoll, QX Jia, B Maiorov, H Wang, and M Maley. Materials science challenges for high-temperature superconducting wire. *Nature materials*, 6(9):631–642, 2007.
- [13] Akihiro Fujii, Koji Shimizu, and Satoshi Watanabe. Efficient exploration of high- T_c superconductors by a gradient-based composition design, 2024. URL <https://arxiv.org/abs/2403.13627>.
- [14] Xin Gui, Bing Lv, and Weiwei Xie. Chemistry in superconductors. *Chemical Reviews*, 121(5):2966–2991, 2021.
- [15] Kam Hamidieh. A data-driven statistical model for predicting the critical temperature of a superconductor. *Computational Materials Science*, 154:346–354, 2018.
- [16] Noah Hoffmann, Tiago FT Cerqueira, Jonathan Schmidt, and Miguel AL Marques. Superconductivity in antiperovskites. *NPJ Computational Materials*, 8(1):150, 2022.
- [17] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [18] Tomohiko Konno, Hodaka Kurokawa, Fuyuki Nabeshima, Yuki Sakishita, Ryo Ogawa, Iwao Hosako, and Atsutaka Maeda. Deep learning model for finding new superconductors. *Physical Review B*, 103(1):014509, 2021.
- [19] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015.
- [20] Shaobo Li, Yabo Dan, Xiang Li, Tiantian Hu, Rongzhi Dong, Zhuo Cao, and Jianjun Hu. Critical temperature prediction of superconductors based on atomic vectors and deep learning. *Symmetry*, 12(2):262, 2020.
- [21] Turab Lookman, Prasanna V Balachandran, Dezhen Xue, and Ruihao Yuan. Active learning in materials science with emphasis on adaptive sampling using uncertainties for targeted design. *npj Computational Materials*, 5(1):21, 2019.

- [22] Aria Mansouri Tehrani, Anton O Oliynyk, Marcus Parry, Zeshan Rizvi, Samantha Couper, Feng Lin, Lowell Miyagi, Taylor D Sparks, and Jakoah Brgoch. Machine learning directed search for ultraincompressible, superhard materials. *Journal of the American Chemical Society*, 140(31): 9844–9853, 2018.
- [23] Efim Mazhnik and Artem R Oganov. Application of machine learning methods for predicting new superhard materials. *Journal of Applied Physics*, 128(7), 2020.
- [24] Bryce Meredig, Erin Antono, Carena Church, Maxwell Hutchinson, Julia Ling, Sean Paradiso, Ben Blaiszik, Ian Foster, Brenna Gibbons, Jason Hattrick-Simpers, et al. Can machine learning identify the next high-temperature superconductor? examining extrapolation performance for materials discovery. *Molecular Systems Design & Engineering*, 3(5):819–825, 2018.
- [25] National Institute of Materials Science. SuperCon dataset, 2022. URL <https://mdr.nims.go.jp/collections/5712mb227>.
- [26] Michael R Norman. The challenge of unconventional superconductivity. *Science*, 332(6026): 196–200, 2011.
- [27] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. Pytorch: An imperative style, high-performance deep learning library. *Advances in neural information processing systems*, 32, 2019.
- [28] Claudio Pereti, Kevin Bernot, Thierry Guizouarn, František Laufek, Anna Vymazalová, Luca Bindi, Roberta Sessoli, and Duccio Fanelli. From individual elements to macroscopic materials: in search of new superconductors via machine learning. *Npj Computational Materials*, 9(1):71, 2023.
- [29] Gordon GC Peterson and Jakoah Brgoch. Materials discovery through machine learning formation energy. *Journal of Physics: Energy*, 3(2):022002, 2021.
- [30] Ghanshyam Pilia, Chenchen Wang, Xun Jiang, Sanguthevar Rajasekaran, and Ramamurthy Ramprasad. Accelerating materials property predictions using machine learning. *Scientific reports*, 3(1):2810, 2013.
- [31] Mario Rodríguez Peña and José Ángel García Guerra. The periodic spiral of elements. *Foundations of Chemistry*, 26(2):315–321, 2024.
- [32] Timo Sommer, Roland Willa, Jörg Schmalian, and Pascal Friederich. 3dsc-a dataset of superconductors including crystal structures. *Scientific Data*, 10(1):816, 2023.
- [33] Valentin Stanev, Corey Oses, A Gilad Kusne, Efrain Rodriguez, Johnpierre Paglione, Stefano Curtarolo, and Ichiro Takeuchi. Machine learning modeling of superconducting critical temperature. *npj Computational Materials*, 4(1):29, 2018.
- [34] Valentin Stanev, Kamal Choudhary, Aaron Gilad Kusne, Johnpierre Paglione, and Ichiro Takeuchi. Artificial intelligence for search and discovery of quantum materials. *Communications Materials*, 2(1):105, 2021.
- [35] Liling Sun and Robert Joseph Cava. High-entropy alloy superconductors: Status, opportunities, and challenges. *Physical Review Materials*, 3(9):090301, 2019.
- [36] Feng Tang, Seishiro Ono, Xiangang Wan, and Haruki Watanabe. High-throughput investigations of topological and nodal superconductors. *Physical Review Letters*, 129(2):027001, 2022.

A Appendix / supplemental material

A.1 Details of Architecture

The convolutional model uses both convolution and pooling with the later stages being fully connected. The fully-connected model obviously uses all to all connections at every layer. The activation function used for both network is ReLU. We provide the visualization of the networks using the Netron software (<https://netron.app/>), based on the saved .onnx files.

A.2 Detailed plots of network performance during training/testing

Below, we present more expansive plots of the training process for accuracy and MSE loss in the two networks as a function of epoch.

A.3 Data leakage

We discuss the possible impact of data leakage [9] on our predictive power. Data leakage happens when the training data contains some kind of information that will not be available when the model is used for prediction or decision making. Such leakage has two main types: target leakage and train-test contamination. Target leakage occurs when training predictors include data that will not be available at the time to make new predictions. That is not relevant for us. As to test-train contamination, there was nothing in our dataset of the form $\text{Mo}_x\text{Re}_y\text{Si}_z$.

If one broadens the definition and considers similar materials, our dataset indeed has materials which share two of the three elements present in the target. While these materials, present in the dataset could be similar, for example, some materials of the form $\text{Mo}_x\text{Re}_y\text{O}_z$, superconductivity is dependent on many subtle effects and is not easy to predict. In fact, different compositions of $\text{Mo}_x\text{Re}_y\text{O}_z$ have different behavior. Also, while we are trying to predict new materials, we expect there are one or more chemical composition-wise similar superconducting materials. Thus, for computationally accessible new superconducting materials, some similarity in the data is expected.

A.4 Details of the experiment

$\text{Mo}_{20}\text{Re}_6\text{Si}_4$ was synthesized successfully by arc melting method, and its crystal structure was confirmed to be tetragonal with lattice parameters $a = b = 9.472(1) \text{ \AA}$ and $c = 4.965(2) \text{ \AA}$. The compound belongs to the σ -phase, which has a large unit cell with 30 atoms and a general stoichiometry of $\text{A}_{20}\text{B}_{10}$ shown in Figure 9a. In binary σ -phases, A and B atoms occupy different Wyckoff positions, and in ternary σ -phases, the third element partially occupies some positions from A and/or B atoms. Although further experiments are required to determine all mixing and site occupancies, assuming the stoichiometry of $\text{Mo}_{20}\text{Re}_6\text{Si}_4$ can be used in further physical properties analysis.

Figure 9b displays the temperature dependence of the zero-field-cooled (ZFC) and field-cooled (FC) volume magnetic susceptibility measured under an applied magnetic field of 20 Oe. The transition to the superconducting state is observed at $T_c = 5.4 \text{ K}$ and was determined as the point of intersection between the steepest slope of χ_v and the extrapolation of the normal state susceptibility. This method is consistent with previous studies. The ZFC volume magnetic susceptibility slightly exceeds $4\pi\chi_v = -1$, which is the expected value for a perfect Meissner effect. Insights into the magnetic properties of the sample and are important for understanding the superconducting behavior of the material.

In order to gain a deeper understanding of the superconducting state of the material, we conducted resistivity measurements. The temperature dependence of the resistivity is depicted in Figure 9c. We observed a decrease in resistivity as the temperature increased, which is typical for metallic behavior. However, the resistivity values were found to be high, likely due to the high level of disorder present in the material. Notably, a sharp drop to zero resistivity was observed at a temperature of approximately 6 K, which is indicative of the superconducting transition.

Finally, to establish the bulk nature of the observed superconductivity, we performed specific heat measurements. The temperature dependence of C/T is presented in Figure 9d. We observed a λ -shaped anomaly, which indicates the presence of a phase transition. To identify the specific heat jump (ΔC) and critical temperature T_c , we constructed an equal-area entropy plot (red lines in the main panel of Figure 9d). The obtained T_c was determined to be 5.4 K, which is in agreement with the magnetic susceptibility data.

A.5 Computational Resources

All networks were trained on a 2 NVIDIA GPU compute nodes managed by the SUN Grid engine, on the Beowulf cluster, Rutgers university. 6 realization of each network were carried out, with the dataset sizes for training enumerated in the main text.

13132 training entries over 5000 epochs required 6 hours and 45 minutes for training for CNN networks and 6 hours and 50 minutes for FC networks. Testing on existing models was completed in 45 seconds (~ 3282 entries).

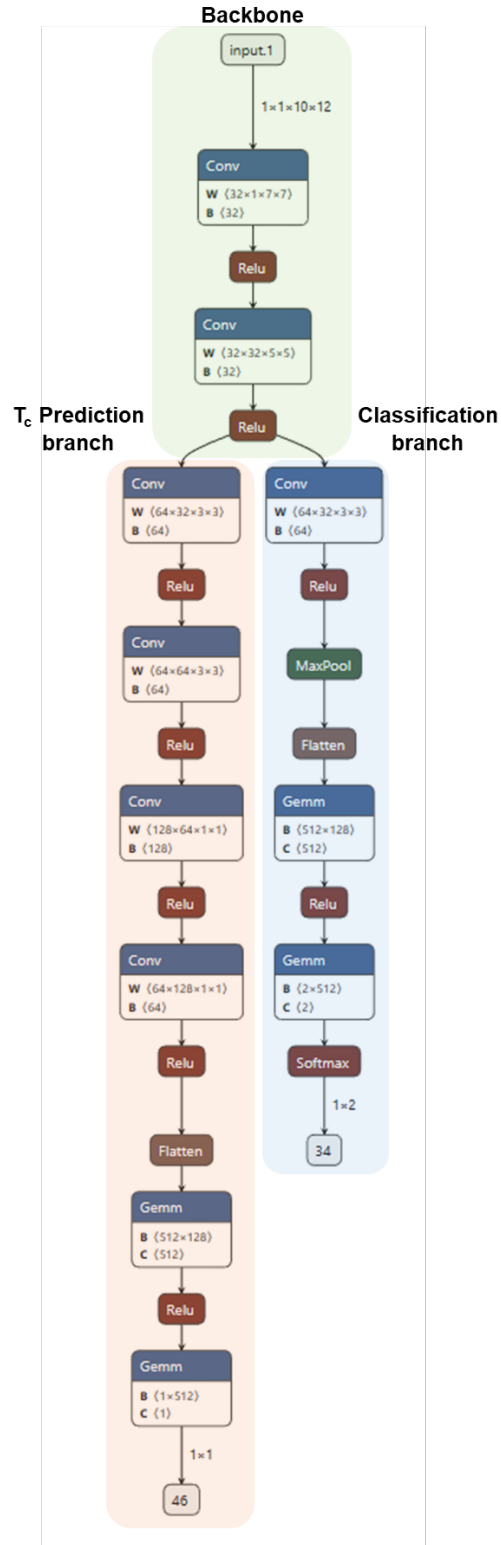


Figure 3: The convolutional model, including (1) Backbone, (2) T_c Prediction branch and (3) Classification branch.

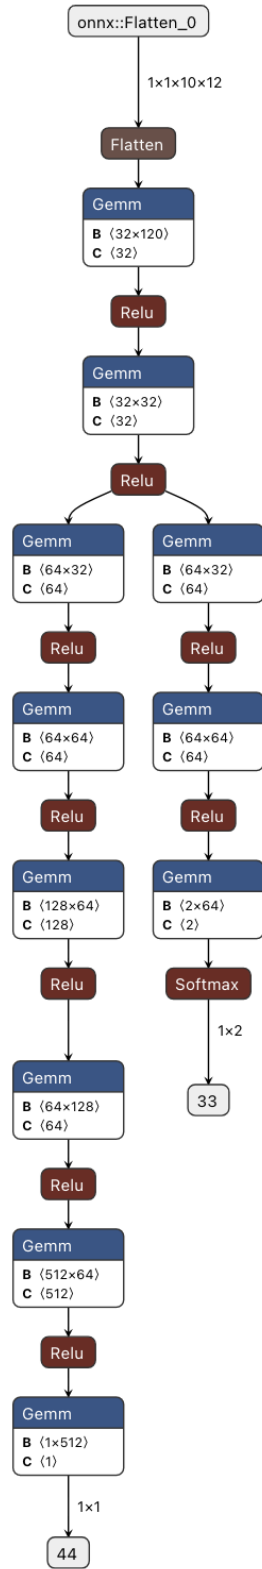


Figure 4: The fully-connected model.




Figure 5: Fully-connected Model: Accuracy in classification during training with epoch.




Figure 6: Fully-connected Model: MSE loss in classification during training.




Figure 7: Convolutional Model: Accuracy in classification during training with epoch.




Figure 8: Convolutional Model: MSE loss in classification during training with epoch.

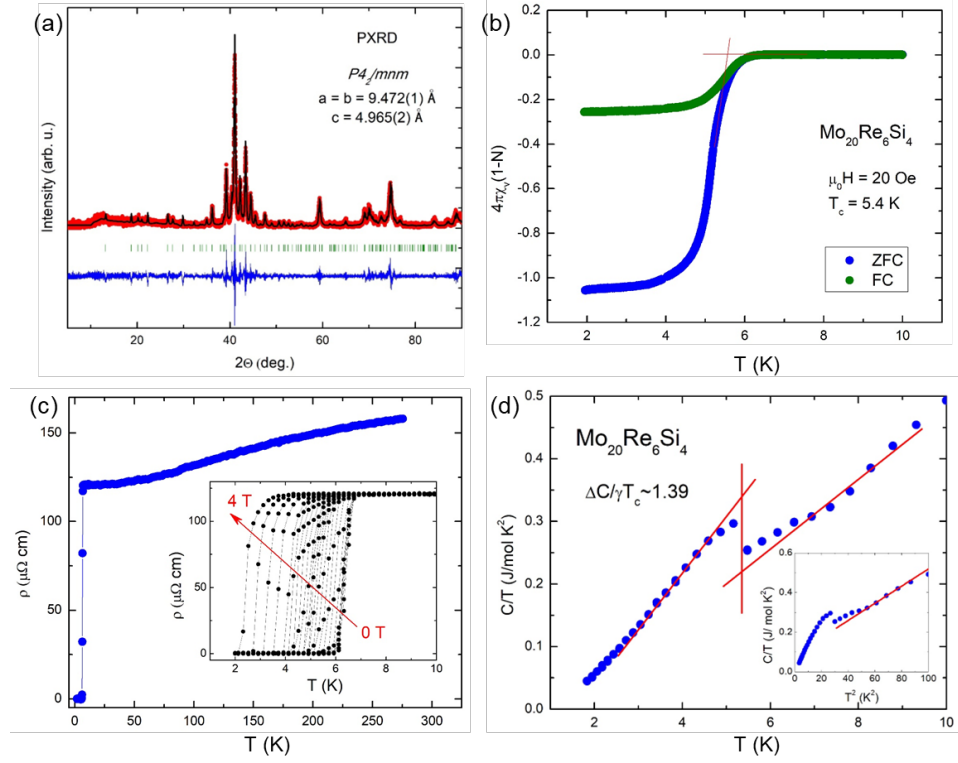


Figure 9: Phase identification and superconductivity properties of $\text{Mo}_{20}\text{Re}_6\text{Si}_4$. (a) Powder X-ray diffraction pattern of sigma phase $\text{Mo}_{20}\text{Re}_6\text{Si}_4$; (b) Magnetic susceptibility of $\text{Mo}_{20}\text{Re}_6\text{Si}_4$. (c) Field-dependent resistivity of $\text{Mo}_{20}\text{Re}_6\text{Si}_4$. (d) Heat capacity of $\text{Mo}_{20}\text{Re}_6\text{Si}_4$.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer , , or .
- means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "" is generally preferable to "", it is perfectly acceptable to answer "" provided a proper justification is given (e.g., "error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering "" or "" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS paper checklist”,**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer:

Justification: Our abstract and the introduction provide an an overview of our computational and experimental results.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer:

Justification: The limitations and possible ways overcome them are discussed in section 5.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer:

Justification: This paper does not have any theoretical results.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer:

Justification: The paper contains statements about both computational studies and physical experiments. Even for the computational part, while sufficient information is given, it will be hard to include it within the 4 page limit. Some more information is available in supplementary materials. If deemed appropriate more information can be made available, if the manuscript is accepted. Similar statements are valid for the physical experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: Yes

Justification: Section 7 discusses data availability and points out the source dataset. We provide the code for this submission in a freely accessible Github repository: <https://github.com/danielkaplan137/TcMLPred.git>

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.

- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer:

Justification: See subsections 2.2, 2.3 and A.1 where most of these details are provided.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer:

Justification: The results in table 1 provide error bars.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer:

Justification: The computational resources are discussed subsection A.5.

Guidelines:

- The answer NA means that the paper does not include experiments.

- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer:

Justification: Since this is a primarily computational work in nature, supported by physical experiments not involving human subjects, which might lead to discovery of new superconductors. This is mostly beneficial for humans if it comes to fruition. We do not think we violate the concerns listed in the Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer:

Justification: While there is some potential future benefits to humanity, the societal impacts of this paper are somewhat limited in scope.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer:

Justification: We expect this work to be used within a limited scope: the initial exploration of possible materials for synthesis. So, these safeguards are not quite appropriate for this work.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer:

Justification: Open-source code and relevant authors are cited when comparing results.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer:

Justification: We are not generating any new asset right away.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer:

Justification: Our work does not involve crowdsourcing and research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer:

Justification: We do not do any research with human subjects and therefore IRB approvals are not relevant.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.