

---

# Simulation-based inference for neutrino interaction model parameter tuning

---

**Karla Tame-Narvaez**

Theory Division

Fermi National Accelerator Laboratory

Batavia, IL 60510

[karla@fnal.gov](mailto:karla@fnal.gov)

**Aleksandra Ćiprijanović**

Computational Science and AI Directorate

Fermi National Accelerator Laboratory

Batavia, IL 60510

Department of Astronomy and Astrophysics

University of Chicago

Chicago, IL 60637

[aleksand@fnal.gov](mailto:aleksand@fnal.gov)

**Steven Gardiner**

Computational Science and AI Directorate

Fermi National Accelerator Laboratory

Batavia, IL 60510

[gardiner@fnal.gov](mailto:gardiner@fnal.gov)

**Giuseppe Cerati**

Computational Science and AI Directorate

Fermi National Accelerator Laboratory

Batavia, IL 60510

[cerati@fnal.gov](mailto:cerati@fnal.gov)

## Abstract

High-energy physics experiments studying neutrinos rely heavily on simulations of their interactions with atomic nuclei. Limitations in the theoretical understanding of these interactions typically necessitate ad hoc tuning of simulation model parameters to data. Traditional tuning methods for neutrino experiments have largely relied on simple algorithms for numerical optimization. While adequate for the modest goals of initial efforts, the complexity of future neutrino tuning campaigns is expected to increase substantially, and new approaches will be needed to make progress. In this paper, we examine the application of simulation-based inference (SBI) to the neutrino interaction model tuning for the first time. Using a previous tuning study performed by the MicroBooNE experiment as a test case, we find that our SBI algorithm can correctly infer the tuned parameter values when confronted with a mock data set generated according to the MicroBooNE procedure. This initial proof-of-principle illustrates a promising new technique for next-generation simulation tuning campaigns for the neutrino experimental community.

## 1 Introduction

A major priority for current research in high-energy physics is the investigation of the properties of neutrinos. Because these elementary particles interact weakly with other forms of matter, typical neutrino experiments involve both intense neutrino sources and large, highly sensitive detectors in order to record enough reactions for analysis. Beyond the technical demands of the required apparatus, however, interpretation of the experimental data poses its own significant difficulties. Among the greatest of these difficulties is the need for precise simulations of collisions between neutrinos and atomic nuclei. Despite ongoing efforts within the scientific community, gaining a full theoretical understanding of these collisions is a formidable problem, and state-of-the-art simulations of neutrino scattering presently rely on many rough, semi-empirical approximations.

To obtain simulation predictions that are reliable enough to use for physics measurements, neutrino experiments have commonly resorted to tuning interaction model parameters to reference data. The widely-used GENIE simulation code [6, 4] has been the most popular platform for such tuning exercises, with several carried out by the developers themselves [27, 28, 29, 19] as well as multiple experimental collaborations [26, 3, 2]. A representative example recently performed by MicroBooNE [2] involved using a software framework called NUISANCE [25] to tune four GENIE model parameters to neutrino scattering data previously obtained by the T2K experiment [1]. A satisfactory adjustment to the base GENIE model, dubbed the “MicroBooNE Tune,” was ultimately achieved via a simple likelihood fit to the data. However, pathological results seen during initial attempts led the authors to take the drastic step of ignoring the reported correlations between bins of the reference T2K measurement. While this stopgap solution, paired with simple numerical methods, was adequate for the modest immediate needs of MicroBooNE, next-generation efforts will be substantially more exacting. To achieve the stringent precision on neutrino interaction modeling required for the goals of the field, future tuning campaigns will necessarily involve both larger parameter spaces and more input data with greater complexity. Technical innovation in simulation tuning procedures provides a potential solution for the neutrino community to overcome these looming obstacles.

Although very recently applied to modeling the performance of the detector hardware in the JUNO neutrino experiment [15], the use of simulation-based inference (SBI) methods for tuning neutrino interaction models has not yet been explored. These methods leverage simulators as a part of the statistical inference procedure, where the goal is to infer the likelihood or posterior distributions for a given experiment or observation. Traditional SBI methods include, for example, Approximate Bayesian Computation [ABC; 24] and Approximate Frequentist Computation [AFC; 11], which are closely related to the traditional template histogram and kernel density estimation approaches. The recent rise of SBI algorithms, which utilize deep neural networks as surrogates for modeling the conditional probability densities of the likelihood or posterior distribution in a given inference problem from simulations, enabled inference even in high-dimensional parameter spaces [14]. Furthermore, after the training of the deep learning model is performed, the subsequent inference is fast and cheap, i.e., the SBI model is amortized.

Deep learning-based SBI models have already been successfully used in other areas of physics. For example, in collider physics [10] for constraining the Higgs potential for di-Higgs production [20], searching for CP violation in leptonic WH production [7], measuring QCD Splittings [8], etc. In astrophysics and cosmology, SBI has been used for the inference of the Hubble constant from binary neutron star mergers [16], the dark matter substructure inference in galaxy-galaxy strong lenses [12, 13, 5], inference of strong lensing parameters [18, 30, 22], inference of galaxy properties from spectra [17], for cosmology inference from galaxy cluster abundance [23], etc.

In this work, we use SBI [9] with a neural posterior estimator (NPE) method to revisit the neutrino interaction model tuning performed by MicroBooNE [2]. SBI provides an efficient alternative to conventional methods, as it offers amortized inference, i.e., after the upfront cost of training the model is paid, inference can be performed in seconds. We demonstrate that our SBI model is capable of inferring correct parameter values when confronted with mock data generated using the MicroBooNE Tune configuration of GENIE. This technical demonstration reveals strong potential for near-future use of SBI to obtain parameter values from actual experimental data sets, thus establishing a novel approach for tuning neutrino interaction simulations.

## 2 Data and Methods

We simulate neutrino-nucleus collisions using GENIE [6, 4]. Within the GENIE interaction model, we vary the four model parameters that were adjusted in the MicroBooNE tune:  $\theta_1$  (MaCCQE),  $\theta_2$  (NormCCMEC),  $\theta_3$  (XSecShape\_CCMEC) and  $\theta_4$  (RPA\_CCQE). The former represent the axial mass for CCQE interactions, which governs the  $Q^2$  dependence of the cross section. NormCCMEC represents the overall normalization of the CC multi-nucleon (MEC) contribution. XSecShape\_CCMEC represent the shape parameter of the CC MEC cross section. Finally, RPA\_CCQE represents which modifies the Random Phase Approximation correction applied to CCQE processes. More detailed physical meanings of these parameters are described in detail in the original MicroBooNE publication [2]. To provide sufficient coverage of the parameter space in the vicinity of the best-fit values from the MicroBooNE tune, we created an ensemble of configurations in which the four parameters were

independently sampled from uniform distributions with the following ranges:

$$\theta_1 \in [0.961, 1.39] \text{ GeV}, \quad \theta_2 \in [1.0, 3.0], \quad \theta_3 \in [0.0, 1.0], \quad \theta_4 \in [0.0, 1.0].$$

For each configuration of the parameters, we processed the output of GENIE with NUISANCE [25] to produce a 58-bin histogram representing a theoretical prediction that can be directly compared to the “Analysis I” T2K neutrino interaction data set reported in Ref. [1]. For the original MicroBooNE tune, this procedure was used together with a subsequent likelihood fit to the measured histogram to obtain best-fit parameter values. In this study, we created a training set of 200,000 configurations and a test set of 1,000 independent configurations to validate our SBI workflow.

The goal of our algorithm is to infer the underlying set of parameters used within GENIE from a prediction histogram generated according to the approach described above. For this task, we utilize the SBI framework developed by `mackelab`<sup>1</sup>, implemented in `python`. The model takes as input the four physics parameters,  $\theta_i$ , where  $i = 1, 2, 3, 4$ , along with their associated histograms,  $x_i$ , to learn the inverse mapping from histograms to parameters.

To facilitate a more efficient learning of the posterior distribution, we use a simple three-layer embedding network to reduce the dimensionality of the histograms from 58 to 24 summary features<sup>2</sup>. These embeddings are then used as inputs to the NPE. The NPE uses a Masked Autoregressive Flow [21] architecture with six transformations and 55 hidden features in each block. Both the embedding network and the MAF are trained together to allow the embedding network to learn the most informative summaries, which will lead to the best posterior predictions.

For the training loop, we use mini-batches of size 512, a learning rate of  $10^{-2}$ , and a training/validation split of 90%/10% to monitor and mitigate overfitting. Training proceeds under an early stopping criterion, with patience of 45 epochs. The training converges in an average of 215 epochs, running in approximately 10 minutes in a regular CPU environment. The code used in this work can be found on our GitHub<sup>3</sup> page.

### 3 Results

The performance of the NPE for an individual randomly chosen test event is illustrated in Figure 1 (top left figure). In here, the diagonal panels show the one-dimensional marginal posterior distributions for each parameter, while the off-diagonal panels display the corresponding two-dimensional joint posteriors. The contours represent the 68% and 95% confidence intervals, enclosing the highest posterior density intervals under the assumption of a well-behaved posterior distribution. The shape and orientation of the contours provide insight into the correlation structure among the parameters: for example, elongated contours indicate strong correlations, while circular contours suggest independence. In this individual test event, we therefore see little correlation between the four parameters; however, some of these correlations will appear when averaged over the whole test sample (for example, Figure 1 top right panel). The vertical and horizontal dashed lines denote the corresponding true parameter values used to generate the plotted test data point. Overall, the posteriors are centered near the target values within the 68% interval.

To assess the overall performance of the model, in Figure 1 (top right) we show the distribution of the residuals  $\Delta\theta_i = \theta_i^{\text{pred}} - \theta_i^{\text{true}}$  for  $i = 1, 2, 3, 4$ , computed over a sample of 1000 independent test events. Again, each diagonal panel displays the one-dimensional distribution of residuals for an individual parameter, while the off-diagonal panels illustrate the joint distributions between pairs of residuals. The dashed lines indicate zero bias. We observe that all residuals are centered around zero with narrow widths, indicating that the model yields unbiased estimates with low variance. The contours in the off-diagonal panels reveal the correlation structure among residuals: mild correlations are visible for all parameter pairs, but no significant systematic deviations are observed.

In Figure 1 (bottom left), we present the posterior coverage probability for each parameter  $\theta_i$ , evaluated again over a sample of 1000 test events. The solid colored lines represent the empirical

---

<sup>1</sup><https://github.com/sbi-dev/sbi>

<sup>2</sup>Comparable results can be obtained with lower-dimensional input data; however, we found that the network becomes overconfident when the number of parameters is reduced below ten. Using 24 parameters proved to be a stable choice, maintaining a well-calibrated model. Further studies on architecture design and optimal parameter tuning are ongoing.

<sup>3</sup>[https://github.com/karlaTame/Neutrino\\_SBI/](https://github.com/karlaTame/Neutrino_SBI/)




Figure 1: Upper left panel: Inferred posterior distributions of four parameters for a single event in the test data set. The gray dashed lines indicate the true values. Upper right panel: Residuals of four parameters for 1000 test events. The gray dashed lines are located at zero for reference. Lower left panel: Posterior coverage of the  $\theta_i$  parameters for 1000 test events. The diagonal black-dashed line indicates perfect uncertainty calibration. The gray regions indicate thresholds of 10% (dark gray) and 20% (light gray) uncertainty miscalibration. Lower right panel: The red points represent the MicroBooNE fit parameters reported in Ref. [2] and used to generate the test histogram  $x_i$ . In blue, we show the four parameters  $\theta_i$  along with their corresponding  $1\sigma$  error bars inferred by our network with  $x_i$  as input. In orange, we show the prior ranges used to train the SBI. We observe an excellent match between the inferred and true parameter values.

coverage for each parameter, indicating the fraction of events where the true value lies within a given posterior confidence interval. The black solid curve represents the combined performance of the model on all parameters, for which we utilize a distance metric  $D(\theta_i)$  from [22].  $D(\theta_i)$  combines parameter values from posterior samples into a single objective function that takes into account the covariance between different  $\theta_i$  parameters. The black dashed diagonal corresponds to perfect uncertainty calibration, where the predicted confidence intervals exactly match the empirical coverage. The shaded bands around the diagonal (dark and light gray) denote 10% and 20% tolerance regions, respectively, for quantifying miscalibration. Two of our inferred parameters are within the 10%, and two within 20% tolerance band. Only  $\theta_3$  shows slight overconfidence, while for the rest of the parameters, the model is slightly underconfident.

Finally, we performed a test to verify the ability of the model to identify the correct parameter values when presented with a histogram generated with the MicroBooNE tune parameters. In Figure 1 (bottom right), we show that the true values are almost identically reproduced. This is a key validation

test towards the application of the model to actual experimental data sets, including the same T2K measurement used to derive the MicroBooNE tune.

## 4 Summary and Conclusion

To overcome present deficiencies in models of neutrino scattering, experimental collaborations often tune simulation parameters to relevant interaction data. Increasingly arduous requirements for these tuning exercises may be addressed by improving the sophistication of the numerical methods applied to them. Here we demonstrate, for the first time, the successful application of advanced AI-based techniques—specifically, simulation-based inference with neural posterior estimation—to efficiently determine the neutrino interaction model parameter values from mock data. Using outputs of the GENIE and NUISANCE codes, we trained an SBI algorithm to infer the correct parameter values when presented with a physics model prediction corresponding to the “MicroBooNE Tune” defined in Ref. [2]. Given that an algorithm trained in this way could be immediately used for parameter inference using actual measurements as input, this successful validation of our approach lays a strong foundation for future neutrino interaction model tuning using SBI techniques.

Future work will apply our SBI algorithm to the T2K measurement studied by MicroBooNE, including a physical interpretation of the inferred parameter values, a full treatment of correlated uncertainties on the inputs and outputs, and an evaluation of the consistency and relative performance of our technique versus the original MicroBooNE likelihood fit. We expect the SBI tuning approach to be easily generalizable to similar problems in neutrino interaction modeling, with the potential to significantly improve the efficiency of future efforts in both time and computing resources.

## Acknowledgments and Disclosure of Funding

This work was produced by FermiForward Discovery Group, LLC under Contract No. 89243024CSC000002 with the U.S. Department of Energy, Office of Science, Office of High Energy Physics. Publisher acknowledges the U.S. Government license to provide public access under the (DOE Public Access Plan). The work of K.T. is supported by DOE Grant KA2401045.

## References

- [1] Ko Abe et al. Measurement of double-differential muon neutrino charged-current interactions on  $C_8H_8$  without pions in the final state using the T2K off-axis beam. *Phys. Rev. D*, 93(11):112012, 2016.
- [2] P. Abratenko et al. New CC $0\pi$  GENIE model tune for MicroBooNE. *Phys. Rev. D*, 105:072001, Apr 2022.
- [3] M. A. Acero et al. Adjusting neutrino interaction models and evaluating uncertainties using NOvA near detector data. *Eur. Phys. J. C*, 80(12):1119, 2020.
- [4] Luis Alvarez-Ruso et al. Recent highlights from GENIE v3. *Eur. Phys. J. ST*, 230(24):4449–4467, 2021.
- [5] Noemi Anau Montel, Adam Coogan, Camila Correa, Konstantin Karchev, and Christoph Weniger. Estimating the warm dark matter mass from strong lensing images with truncated marginal neural ratio estimation. *Mon. Not. R. Astron. Soc.*, 518(2):2746–2760, January 2023.
- [6] C. Andreopoulos et al. The GENIE neutrino Monte Carlo generator. *Nucl. Instrum. Methods Phys. Res. A*, 614(1):87–104, 2010.
- [7] Ricardo Barrué, Patricia Conde Muñó, Valerio Dao, and Rui Santos. Simulation-based inference in the search for CP violation in leptonic WH production. *Journal of High Energy Physics*, 2024(4):14, April 2024.
- [8] Sebastian Bieringer, Anja Butter, Theo Heimel, Stefan Höche, Ullrich Köthe, Tilman Plehn, and Stefan T. Radev. Measuring QCD Splittings with Invertible Networks. *SciPost Physics*, 10(6):126, June 2021.

- [9] Jan Boelts, Michael Deistler, Manuel Gloeckler, Álvaro Tejero-Cantero, Jan-Matthis Lueckmann, Guy Moss, Peter Steinbach, Thomas Moreau, Fabio Muratore, Julia Linhart, Conor Durkan, Julius Vetter, Benjamin Kurt Miller, Maternus Herold, Abolfazl Ziaeemehr, Matthijs Pals, Theo Gruner, Sebastian Bischoff, Nastya Krouglova, Richard Gao, Janne K. Lappalainen, Bálint Mucsányi, Felix Pei, Auguste Schulz, Zinovia Stefanidi, Pedro Rodrigues, Cornelius Schröder, Faried Abu Zaid, Jonas Beck, Jaivardhan Kapoor, David S. Greenberg, Pedro J. Gonçalves, and Jakob H. Macke. sbi reloaded: a toolkit for simulation-based inference workflows. *Journal of Open Source Software*, 10(108):7754, 2025.
- [10] Johann Brehmer. Simulation-based inference in particle physics. *Nature Reviews Physics*, 3(5):305–305, May 2021.
- [11] Johann Brehmer, Kyle Cranmer, Gilles Louppe, and Juan Pavez. A guide to constraining effective field theories with machine learning. *Phys. Rev. D*, 98(5), September 2018.
- [12] Johann Brehmer, Siddharth Mishra-Sharma, Joeri Hermans, Gilles Louppe, and Kyle Cranmer. Mining for dark matter substructure: Inferring subhalo population properties from strong lenses with machine learning. *Astrophys. J.*, 886(1):49, nov 2019.
- [13] Adam Coogan, Noemi Anau Montel, Konstantin Karchev, Meiert W. Grootes, Francesco Nattino, and Christoph Weniger. One never walks alone: the effect of the perturber population on subhalo measurements in strong gravitational lenses. *arXiv e-prints*, page arXiv:2209.09918, September 2022.
- [14] Kyle Cranmer, Johann Brehmer, and Gilles Louppe. The frontier of simulation-based inference. *Proceedings of the National Academy of Science*, 117(48):30055–30062, December 2020.
- [15] A. Gavrikov, A. Serafini, D. Dolzhikov, A. Garfagnini, M. Gonchar, M. Grassi, R. Brugnera, V. Cerrone, L. V. D'Auria, R. M. Guizzetti, L. Lastrucci, G. Andronico, V. Antonelli, A. Barresi, D. Basilico, M. Beretta, A. Bergnoli, M. Borghesi, A. Brigatti, R. Bruno, A. Budano, B. Caccianiga, A. Cammi, R. Caruso, D. Chiesa, C. Clementi, C. Coletta, S. Dusini, A. Fabbri, G. Felici, G. Ferrante, M. G. Giannmarchi, N. Giudice, N. Guardone, F. Houria, C. Landini, I. Lippi, L. Loi, P. Lombardi, F. Mantovani, S. M. Mari, A. Martini, L. Miramonti, M. Montuschi, M. Nastasi, D. Orestano, F. Ortica, A. Paoloni, L. Pelicci, E. Percalli, F. Petrucci, E. Previtali, G. Ranucci, A. C. Re, B. Ricci, A. Romani, C. Sirignano, M. Sisti, L. Stanco, E. Stanescu Farilla, V. Strati, M. D. C Torri, C. Tuvè, C. Venettacci, G. Verde, and L. Votano. Simulation-based inference for precision neutrino physics through neural monte carlo tuning, 2025.
- [16] Francesca Gerardi, Stephen M. Feeney, and Justin Alsing. Unbiased likelihood-free inference of the Hubble constant from light standard sirens. *Phys. Rev. D*, 104(8):083531, October 2021.
- [17] Gourav Khullar, Brian Nord, Aleksandra Ćiprijanović, Jason Poh, and Fei Xu. DIGS: deep inference of galaxy spectra with neural posterior estimation. *Machine Learning: Science and Technology*, 3(4):04LT04, December 2022.
- [18] Ronan Legin, Yashar Hezaveh, Laurence Perreault Levasseur, and Benjamin Wandelt. Simulation-Based Inference of Strong Gravitational Lensing Parameters. *arXiv e-prints*, page arXiv:2112.05278, December 2021.
- [19] Weijun Li et al. First combined tuning on transverse kinematic imbalance data with and without pion production constraints. *Phys. Rev. D*, 110(7):072016, 2024.
- [20] Radha Mastandrea, Benjamin Nachman, and Tilman Plehn. Constraining the Higgs potential with neural simulation-based inference for di-Higgs production. *Phys. Rev. D*, 110(5):056004, September 2024.
- [21] George Papamakarios, Theo Pavlakou, and Iain Murray. Masked autoregressive flow for density estimation, 2018.
- [22] Jason Poh, Ashwin Samudre, Aleksandra Ćiprijanović, Joshua Frieman, Gourav Khullar, and Brian D. Nord. Deep inference of simulated strong lenses in ground-based surveys. *J. Cosmol. Astropart. Phys.*, 2025(5):053, May 2025.

- [23] Moonzarin Reza, Yuanyuan Zhang, Brian Nord, Jason Poh, Aleksandra Ciprijanovic, and Louis Strigari. Estimating Cosmological Constraints from Galaxy Cluster Abundance using Simulation-Based Inference. In *Machine Learning for Astrophysics*, page 20, July 2022.
- [24] Donald B. Rubin. Bayesianly justifiable and relevant frequency calculations for the applied statistician. *The Annals of Statistics*, 12(4):1151–1172, 1984.
- [25] P. Stowell et al. NUISANCE: a neutrino cross-section generator tuning and comparison framework. *JINST*, 12(01):P01016, 2017.
- [26] P. Stowell et al. Tuning the GENIE Pion Production Model with MINER $\nu$ A Data. *Phys. Rev. D*, 100(7):072005, 2019.
- [27] Júlia Tena-Vidal et al. Neutrino-nucleon cross-section model tuning in GENIE v3. *Phys. Rev. D*, 104(7):072009, 2021.
- [28] Júlia Tena-Vidal et al. Hadronization model tuning in genie v3. *Phys. Rev. D*, 105(1):012009, 2022.
- [29] Julia Tena-Vidal et al. Neutrino-nucleus CC0 $\pi$  cross-section tuning in GENIE v3. *Phys. Rev. D*, 106(11):112001, 2022.
- [30] Sebastian Wagner-Carena, Jelle Aalbers, Simon Birrer, Ethan O. Nadler, Elise Darragh-Ford, Philip J. Marshall, and Risa H. Wechsler. From Images to Dark Matter: End-to-end Inference of Substructure from Hundreds of Strong Gravitational Lenses. *Astrophys. J.*, 942(2):75, January 2023.