
The Pareto frontier of resilient jet tagging

Anonymous Author(s)

Affiliation
Address
email

Abstract

1 Classifying hadronic jets using their constituents' kinematic information is a critical
2 task in modern high-energy collider physics. Often, classifiers are designed by
3 targeting the best performance using metrics such as accuracy, AUC, or rejection
4 rates. However, the use of a single metric can lead to the use of architectures
5 that are more model-dependent than competitive alternatives, leading to potential
6 uncertainty and bias in analysis. We explore such trade-offs and demonstrate the
7 consequences of using networks with high performance metrics but low resilience.

8 **1 Introduction**

9 When strongly-interacting quarks and gluons are produced by high-energy particle collisions at
10 colliders like the Large Hadron Collider (LHC), they shower and hadronize, creating a collimated
11 ‘jet’ of particles in the final state that is imprinted with some properties of the originating particle (1).
12 Classification, or *tagging*, of these jets based on their *substructure* has become a critical task at the
13 Large Hadron Collider (LHC), where many studies require doing so to extract maximal information
14 from the data (2). Jet tagging has become the quintessential proving grounds for Artificial Intelligence
15 / Machine Learning (AI/ML) algorithms at the LHC: state-of-the-art transformer and graph-based
16 architectures (3; 4; 5; 6; 7) are significantly more performant than earlier approaches (8; 9; 10; 11; 12).
17 Economists say, “When a measure becomes a target, it ceases to be a good measure.” (13) While
18 the accuracy of an ML/AI classifier, often measured by ‘AUC’ (area under the ROC curve), is a
19 critical benchmark, fixation on this quantity can lead to sub-optimal outcomes in analyses. As model
20 complexity increases, they can become susceptible to learning idiosyncrasies of the simulated training
21 sample rather than genuine generalizable physics information. This has been studied by ATLAS (14),
22 which showed that classifiers are more susceptible to uncertainties related to physics modeling than
23 those related to detector effects. Similar studies have explored solutions to generalizability (15).
24 In this work, we evaluate architectures that are often used for tagging in terms of their AUC and their
25 simulation model-dependence, or ‘resilience’ (16; 17). Models with varying complexities were trained
26 and tested on *different* Monte Carlo (MC) simulated datasets; then used to construct the ‘Pareto
27 frontier’ (18) of AUC vs. resilience. We perform a case study to demonstrate the risk of biasing
28 downstream parameter estimation tasks when using models with low resilience. We advocate for a
29 holistic approach to classifier design that includes multiple benchmarks, suited to the application.

30 **2 Methodology**

31 **2.1 Monte Carlo simulated event samples**

32 Two jet classification tasks were considered in these studies: the discrimination of jets initiated by
33 a quark or gluon (‘q/g tagging’), and the identification of jets resulting from the hadronic decay of
34 a Lorentz-boosted top quark (‘top tagging’). For each of these tasks, a set of MC simulated events

35 generated with PYTHIA 8 (19) was used to train the classifiers in a fully-supervised manner. PYTHIA
 36 samples used the default Monash set of tuned parameters (20) in all cases. Alternative samples of the
 37 same processes were also generated with HERWIG 7 (21)¹, to enable quantification of the resilience
 38 as the AUC %-difference between testing on the nominal and alternative sample.² All jets, regardless
 39 of their size, are reconstructed with FASTJET (22) and filtered to have a transverse momentum (p_T)
 40 between 500–550 GeV. No detector simulation is applied.

41 For q/g tagging studies, the simulated event samples from Refs. (23) were used, which are freely
 42 available on the CERN Zenodo platform (24; 25). These samples consist of one million anti- k_t
 43 $R = 0.4$ jets from each of the $Z + q$ or $Z + g$ processes, where the Z boson decays into neutrinos.
 44 The boosted top tagging studies were performed using a new set of samples, which were generated for
 45 these studies and made publicly available on the CERN Zenodo platform (26). One million mixed q/g
 46 background jets for this task with the same kinematic selection were produced from a dijet process.

47 2.2 Model architectures surveyed

48 We have surveyed a representative selection of architectures that are either currently used or have
 49 recently been used in physics analysis at the LHC, and summarize the setups studied in Table 1.
 50 All networks were trained in a fully-supervised manner for 500 epochs³, using early stopping and
 51 a patience of 10 epochs. The default ADAM (27) optimizer was used, with a learning rate of 0.001.
 52 The MC samples were split such that 75% of the events were used for training, and 12.5% each were
 53 used for testing & validation. Each network was given only particle-level kinematic information
 54 (constituent p_T , pseudorapidity η , and azimuthal angle ϕ) as input.

Table 1: An overview of the jet tagging architectures surveyed in this study. For each tagger, the key hyperparameters that were scanned to explore the performance-resilience trade-off are listed.

Architecture	Hyperparameters Scanned	Reference
Expert Features	Angularities: β values Multiplicities: p_T cuts, charge req.	(28)
Deep Neural Networks (DNNs)	Hidden layers: 2–10 Neurons per layer: 1–300	(23)
Particle-Flow Networks (PFNs)	Latent dimension ℓ : 1–1024 Φ network nodes: 50–500	(23)
Energy-Flow Networks (EFNs)	Same as PFNs	(23)
Particle Transformer (ParT)	Attention heads: 2, 4, 8	(29)

55 3 Results

56 3.1 Pareto Frontier

57 Figure 1 shows the classifier AUC vs. resilience (AUC %-difference for Pythia vs. Herwig samples),
 58 for each of the models trained and for each set of hyperparameters considered: ‘optimal’ performance
 59 is at the lower-right corner of the figure. The Pareto frontier connecting models that optimize the
 60 AUC-resiliency tradeoff is highlighted, and models in the shaded region are Pareto-excluded.

61 The Pareto frontier shows that more “complex” models (e.g. ParT) do achieve a higher raw perfor-
 62 mance, but at the cost of resiliency. On the other hand, simpler models based on physical principles,
 63 such as EFNs or mere expert features like angularities, are more robust. For top tagging, vertical
 64 columns of different network types strongly discourage the use of unnecessarily complex networks.

¹For q/g (top), samples are generated with PYTHIA version 8.226 (8.331) and HERWIG version 7.1.4 (7.3.0).

²There are many potential ways to define resilience — just as AUC does not capture all important aspects of classification, so too does this particular definition of resilience.

³This number was chosen arbitrarily, the early stopping condition causes training to terminate in between 30–100 epochs in most cases.

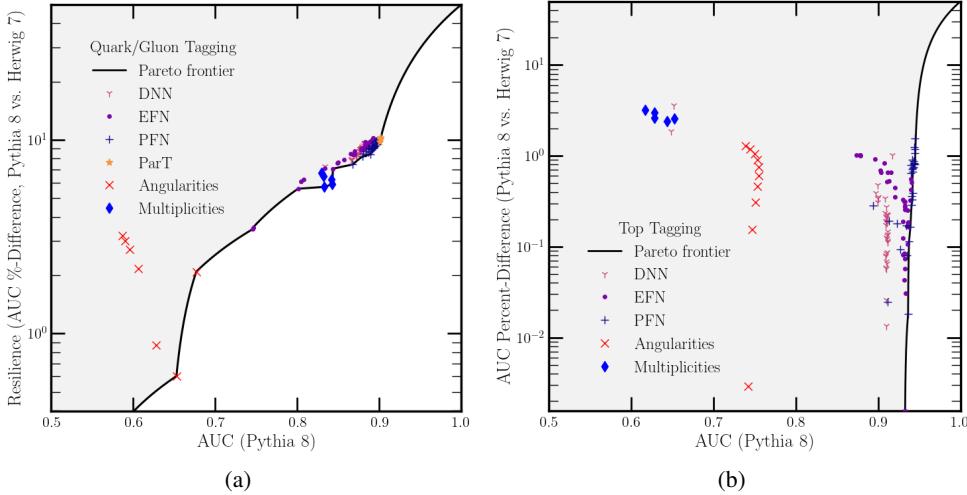


Figure 1: The Pareto frontier for (a) q/g tagging, (b) top tagging tasks. The AUC of models trained with PYTHIA samples is plotted *vs.* the resilience, defined as the percent difference in AUC evaluated on the PYTHIA and HERWIG test samples. The shaded grey region is Pareto-excluded.

65 3.2 Knowledge Distillation

66 In an attempt to overcome the Pareto frontier and make models better along both AUC and resiliency,
 67 we tried Knowledge Distillation: a complex ‘teacher’ model is used to train a less complex ‘student’
 68 model (30). Ultimately, this approach was unsuccessful in overcoming the Pareto frontier, but
 69 interesting observations made during the study motivate us to document it here. The PFN with
 70 $\ell = 128$ and 250 dense nodes per hidden layer was used as the teacher for this study, while various
 71 DNN and EFN models were used as students. The training procedure in Section 2.2 was modified
 72 such that the students were trained instead using the teacher’s prediction as ‘soft labels’ by minimizing
 73 the KL-divergence between the predictions of the teacher and student models per-batch. Various
 74 forms of regularization were tested, but no significant change in the outcome was found.

75 The results for a representative pair of student models are shown in the AUC-Resilience plane in
 76 Figure 2a, along with the teacher model and ‘baseline’ models whose architectures match the students,
 77 but that are not trained using distillation. The contour between the baseline and teacher models is
 78 also drawn: it is obtained by performing inference with a linear combination of the two models on
 79 the test set that varies from pure-teacher to pure-baseline. The student models beat this contour,
 80 demonstrating non-trivial improvement: the AUC of the model increases more than its resilience
 81 degrades. However, when models that are closer to the Pareto frontier are selected for distillation, the
 82 observed improvement is reduced. The results of distilling into the many DNN and EFN models we
 83 study is shown in Figure 2b: while many students improve, none surpass the existing frontier.

84 3.3 Case Study: determining q/g fractions

85 For a realistic downstream analysis task, a less accurate but more resilient classifier can ultimately
 86 yield a less biased physics result. We illustrate this with a case study, where the flavor *mixture fraction*
 87 κ of a mixed sample of quark and gluon jets is estimated using two PFNs located on the Pareto
 88 frontier: a small, resilient PFN ($\ell = 8$, 50 nodes per hidden layer) and a large network with a higher
 89 AUC ($\ell = 128$, 250 nodes per hidden layer). Given a classifier, one can extract the per-event q/g
 90 likelihood ratio $\ell(x)$ via the Neyman Pearson Lemma (31), from which the κ maximum likelihood
 91 inferences may be extracted. The flavor composition of two mixed jet samples with respective
 92 quark-initiated jet fractions of 50% and 25% is estimated with both networks, for jets from either the
 93 PYTHIA or HERWIG samples; the results are tabulated in Table 2. The experiment is re-run 10 times
 94 with re-trained networks as a proxy for confidence interval estimation.

95 Both the large and small networks are able to accurately recover the mixture fraction when samples
 96 are constructed with PYTHIA jets. However, when classifiers trained with PYTHIA jets are used
 97 to estimate the mixture fraction in a sample of HERWIG jets (used as pseudodata sample in this

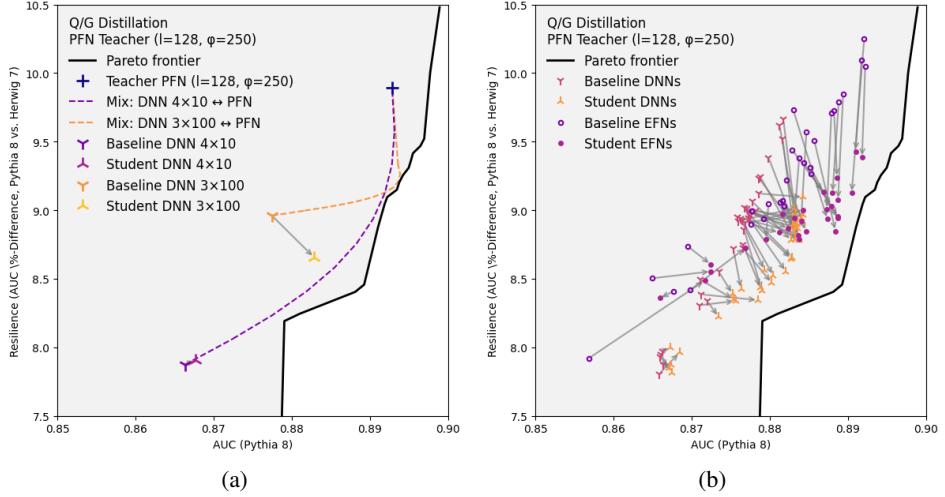


Figure 2: (a) Results of training two student DNNs via distillation from a teacher PFN. (b) Summary of distillation training from a teacher PFN to all DNNs and EFNs in the study.

study), the inferred κ value is biased. The extraction process can be calibrated by reweighting using a second set of classifiers that distinguish between PYTHIA and HERWIG. This classifier approximates the likelihood ratio between the two classes, allowing one to reweight one class of samples to be statistically identical from the other class (modulo reweighting uncertainties)⁴. The PFN models used for the reweighting are identical in architecture to those used for classification. Following calibration, we see from Table 2 that the less resilient model is still biased: the inferred $\hat{\kappa}$ values are not statistically consistent with the true κ values. The more resilient model is unbiased (within 2σ) following the calibration procedure, despite its naively worse performance according to the AUC. The conclusions of this study may apply broadly: any perturbation of the correlation structure of a sample with respect to the training set (*e.g.* fast vs. full detector simulation) may result in such a bias for parameter extractions depending on a unresilient classifier. This is particularly relevant to substructure, as predictions are known to differ from each other and from measurement (33; 34; 35; 36; 37; 38; 39; 40).

Table 2: Summary of the q/g mixture fraction (κ) estimation case study. Reported uncertainties are determined from the standard error from multiple pseudo-experiments.

Classifier	Pythia 8		Herwig 7		Result
	True κ	Inferred $E[\hat{\kappa}]$	Inferred $E[\hat{\kappa}]$	Calibrated $E[\hat{\kappa}]$	
Large PFN	0.50	0.490 ± 0.005	0.296 ± 0.007	0.529 ± 0.006	Biased ✗
	0.25	0.253 ± 0.005	0.125 ± 0.005	0.305 ± 0.006	Biased ✗
Small PFN	0.50	0.504 ± 0.013	0.336 ± 0.016	0.478 ± 0.017	Unbiased ✓
	0.25	0.258 ± 0.013	0.157 ± 0.014	0.268 ± 0.013	Unbiased ✓

4 Concluding remarks

There is a clear trade-off between classifier performance and resilience, which we have visualized in these studies as a Pareto frontier. We have found that the complexity of a given model is a primary driver along the frontier, and that suboptimal model architectures can be improved with more sophisticated approaches to training such as knowledge distillation. Ultimately, the choice of a classifier that is not resilient can lead to suboptimal performance and increased bias in downstream tasks, even if the model is more accurate than others, motivating a more holistic approach to classifier development that includes multiple benchmarks.

⁴In principle, one can extract reweighting uncertainties and obtain confidence intervals using a method such as WiFi ensembles (32), but for simplicity we do not do this here.

118 **References**

- 119 [1] G. P. Salam, “Towards Jetography,” *Eur. Phys. J. C*, vol. 67, pp. 637–686, 2010.
- 120 [2] S. Marzani, G. Soyez, and M. Spannowsky, *Looking inside jets: an introduction to jet substructure and*
121 *boosted-object phenomenology*, vol. 958. Springer, 2019.
- 122 [3] ATLAS Collaboration, “Constituent-Based Quark Gluon Tagging using Transformers with the ATLAS
123 detector.” ATL-PHYS-PUB-2023-032, 2023.
- 124 [4] ATLAS Collaboration, “Constituent-Based W -boson Tagging with the ATLAS Detector.” ATL-PHYS-
125 PUB-2023-020, 2023.
- 126 [5] ATLAS Collaboration, “Tagging boosted W bosons applying machine learning to the Lund Jet Plane.”
127 ATL-PHYS-PUB-2023-017, 2023.
- 128 [6] ATLAS Collaboration, “Constituent-Based Top-Quark Tagging with the ATLAS Detector.” ATL-PHYS-
129 PUB-2022-039, 2022.
- 130 [7] CMS Collaboration, “Identification of heavy, energetic, hadronically decaying particles using machine-
131 learning techniques,” *JINST*, vol. 15, p. P06005, 2020.
- 132 [8] CMS Collaboration, “Identification techniques for highly boosted W bosons that decay into hadrons,”
133 *JHEP*, vol. 12, p. 017, 2014.
- 134 [9] ATLAS Collaboration, “Identification of boosted, hadronically decaying W bosons and comparisons with
135 ATLAS data taken at $\sqrt{s} = 8$ TeV,” *Eur. Phys. J. C*, vol. 76, p. 154, 2016.
- 136 [10] ATLAS Collaboration, “Identification of high transverse momentum top quarks in pp collisions at $\sqrt{s} =$
137 8 TeV with the ATLAS detector,” *JHEP*, vol. 06, p. 093, 2016.
- 138 [11] ATLAS Collaboration, “Performance of top-quark and W -boson tagging with ATLAS in Run 2 of the
139 LHC,” *Eur. Phys. J. C*, vol. 79, p. 375, 2019.
- 140 [12] ATLAS Collaboration, “Performance and calibration of quark/gluon-jet taggers using 140 fb^{-1} of pp
141 collisions at $\sqrt{s} = 13$ TeV with the ATLAS detector,” 2023.
- 142 [13] M. J. Artis, “Monetary Theory and Practice. The UK Experience,” *The Economic Journal*, vol. 94,
143 pp. 984–985, 12 1984.
- 144 [14] ATLAS Collaboration, “Accuracy versus precision in boosted top tagging with the ATLAS detector,”
145 *JINST*, vol. 19, no. 08, p. P08018, 2024.
- 146 [15] A. Butter, B. M. Dillon, T. Plehn, and L. Vogel, “Performance versus resilience in modern quark-gluon
147 tagging,” *SciPost Phys. Core*, vol. 6, p. 085, 2023.
- 148 [16] G. Soyez, “Pileup mitigation at the LHC: A theorist’s view,” *Phys. Rept.*, vol. 803, pp. 1–158, 2019.
- 149 [17] *Les Houches 2017: Physics at TeV Colliders Standard Model Working Group Report*, 3 2018.
- 150 [18] E. Zitzler, J. Knowles, and L. Thiele, *Quality Assessment of Pareto Set Approximations*, p. 373–404. Berlin,
151 Heidelberg: Springer-Verlag, 2008.
- 152 [19] T. Sjöstrand, S. Mrenna, and P. Skands, “A brief introduction to PYTHIA 8.1,” *Comput. Phys. Commun.*,
153 vol. 178, pp. 852–867, 2008.
- 154 [20] P. Skands, S. Carrazza, and J. Rojo, “Tuning PYTHIA 8.1: the Monash 2013 Tune,” *Eur. Phys. J. C*, vol. 74,
155 no. 8, p. 3024, 2014.
- 156 [21] J. Bellm *et al.*, “Herwig 7.1 Release Note,” 2017.
- 157 [22] M. Cacciari, G. P. Salam, and G. Soyez, “FastJet User Manual,” *Eur. Phys. J. C*, vol. 72, p. 1896, 2012.
- 158 [23] P. T. Komiske, E. M. Metodiev, and J. Thaler, “Energy Flow Networks: Deep Sets for Particle Jets,” *JHEP*,
159 vol. 01, p. 121, 2019.
- 160 [24] P. T. Komiske, E. M. Metodiev, and J. Thaler, “Pythia8 quark and gluon jets for energy flow, v1.” Zenodo,
161 May 2019. <https://doi.org/10.5281/zenodo.3164691>.
- 162 [25] A. Pathak, P. T. Komiske, E. M. Metodiev, and M. Schwartz, “Herwig7.1 quark and gluon jets, v1.” Zenodo,
163 May 2019. <https://doi.org/10.5281/zenodo.3066475>.

- 164 [26] “Pythia8 and herwig7 boosted top & qcd jet datasets.” Online, 8 2025.
 165 <https://cernbox.cern.ch/s/ZyL5WzrlDZmLtuq>.
- 166 [27] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” 2017.
- 167 [28] A. J. Larkoski, J. Thaler, and W. J. Waalewijn, “Gaining (Mutual) Information about Quark/Gluon
 168 Discrimination,” *JHEP*, vol. 11, p. 129, 2014.
- 169 [29] H. Qu, C. Li, and S. Qian, “Particle transformer for jet tagging,” in *Proceedings of the 39th International
 170 Conference on Machine Learning* (K. Chaudhuri, S. Jegelka, L. Song, C. Szepesvari, G. Niu, and S. Sabato,
 171 eds.), vol. 162 of *Proceedings of Machine Learning Research*, pp. 18281–18292, PMLR, 17–23 Jul 2022.
- 172 [30] G. Hinton, O. Vinyals, and J. Dean, “Distilling the knowledge in a neural network,” 2015.
- 173 [31] J. Neyman and E. S. Pearson, “On the problem of the most efficient tests of statistical hypotheses,”
 174 *Philosophical Transactions of the Royal Society of London. Series A, Containing Papers of a Mathematical
 175 or Physical Character*, vol. 231, pp. 289–337, 1933.
- 176 [32] S. Benevides and J. Thaler, “Frequentist Uncertainties on Neural Density Ratios with wifi Ensembles,” 5
 177 2025.
- 178 [33] ATLAS Collaboration, “Measurement of the Soft-Drop Jet Mass in pp Collisions at $\sqrt{s} = 13$ TeV with
 179 the ATLAS detector,” *Phys. Rev. Lett.*, vol. 121, p. 092001, 2018.
- 180 [34] ATLAS Collaboration, “Measurement of soft-drop jet observables in pp collisions with the ATLAS detector
 181 at $\sqrt{s} = 13$ TeV,” *Phys. Rev. D*, vol. 101, p. 052007, 2020.
- 182 [35] ATLAS Collaboration, “Properties of jet fragmentation using charged particles measured with the ATLAS
 183 detector in pp collisions at $\sqrt{s} = 13$ TeV,” *Phys. Rev. D*, vol. 100, p. 052011, 2019.
- 184 [36] ATLAS Collaboration, “Measurement of the Lund Jet Plane Using Charged Particles in 13 TeV Proton–
 185 Proton Collisions with the ATLAS Detector,” *Phys. Rev. Lett.*, vol. 124, p. 222002, 2020.
- 186 [37] ATLAS Collaboration, “Measurements of Lund subjet multiplicities in 13 TeV proton-proton collisions
 187 with the ATLAS detector,” *Phys. Lett. B*, vol. 859, p. 139090, 2024.
- 188 [38] ATLAS Collaboration, “Measurement of jet track functions in pp collisions at $s=13$ TeV with the ATLAS
 189 detector,” *Phys. Lett. B*, vol. 868, p. 139680, 2025.
- 190 [39] CMS Collaboration, “Study of quark and gluon jet substructure in $Z+jet$ and dijet events from pp collisions,”
 191 *JHEP*, vol. 01, p. 188, 2022.
- 192 [40] CMS Collaboration, “Measurement of the primary Lund jet plane density in proton-proton collisions at \sqrt{s}
 193 = 13 TeV,” *JHEP*, vol. 05, p. 116, 2024.