

---

# Diffusion Autoencoders with Perceivers for Long, Irregular and Multimodal Astronomical Sequences

---

**Anonymous Author(s)**

Affiliation

Address

email

## 1 1 Introduction

2 Astronomical datasets continue to swell in volume and complexity, driven by large-scale surveys.  
3 Beyond the cataloging of galaxies and stars – photometrically by SDSS [York et al., 2000] and  
4 spectroscopically by DESI [Abareshi et al., 2022] – imaging surveys such as the Asteroid Terrestrial  
5 Last Alert System [ATLAS; Tonry et al., 2018], the Zwicky Transient Facility [ZTF; Bellm et al.,  
6 2018, Masci et al., 2018, Graham et al., 2019, Dekany et al., 2020], and the Young Supernovae  
7 Experiment [YSE; Jones et al., 2021, Aleo et al., 2023] now deliver nearly continuous coverage of  
8 the night sky in search of temporally evolving phenomena.

9 To extract meaningful structure from these massive datasets, self-supervised learning (SSL) has  
10 emerged as a powerful paradigm. SSL objectives are now widely used to pre-train large astrophysical  
11 networks [Zhang et al., 2024, Rizhko and Bloom, 2025], improving robustness to out-of-distribution  
12 data and enabling superior performance when fine-tuned on downstream tasks [Rizhko and Bloom,  
13 2025].

14 A common SSL approach is reconstruction, in which the model learns to reproduce its input. Score-  
15 based generative models have achieved high-fidelity reconstructions in the image domain [Ho et al.,  
16 2020, Dhariwal and Nichol, 2021], but their latent spaces typically lack alignment with high-level  
17 semantic features. To address this limitation, Preechakul et al. [2022] introduced the diffusion  
18 autoencoder, which combines an autoencoder for dimensionality reduction with a diffusion model  
19 conditioned on the encoder latent space. This hybrid approach produces both semantically meaningful  
20 features and reconstructions of exceptional quality.

21 While diffusion autoencoders have shown promise for feature learning on image data, they are not  
22 directly applicable to astronomical time-series data. Time-domain astrophysics, in particular, involves  
23 long sequences of irregularly sampled, noisy photometric and spectroscopic data. These datasets  
24 demand custom architectures for scalable representation learning.

## 25 2 Background

26 **Diffusion autoencoders.** Originally proposed for images by Preechakul et al. [2022], diffusion  
27 autoencoders encode data with an encoder and reconstruct them using a conditional diffusion model.  
28 Since the encoding guides every denoising step, they can capture fine details more effectively than,  
29 for example, variational autoencoders [Kingma and Welling, 2013]. However, the original diffusion  
30 autoencoder relied on U-Net [Preechakul et al., 2022, Dhariwal and Nichol, 2021], which is better  
31 suited to regular modalities like images.

32 **Perceiver.** Perceiver and Perceiver-IO [Jaegle et al., 2021b,a] provide a general framework to (1)  
33 encode irregularly sampled sequences into a latent representation and (2) query outputs from this  
34 latent. This makes them a natural fit for integration with diffusion transformers [Dhariwal and Nichol,  
35 2021], enabling scalable representation learning with diffusion autoencoders.

36 **Related work.** Autoencoding and dimensionality reduction have a long history in representation  
 37 learning. Early models that remain widely used include variational autoencoders [VAEs, Kingma  
 38 and Welling, 2013] and their variants, such as hierarchical models [Vahdat and Kautz, 2020] and  
 39 discrete latent spaces [Van Den Oord et al., 2017, Razavi et al., 2019]. These models remain common  
 40 in physics applications, though they often suffer from posterior collapse [Van Den Oord et al., 2017,  
 41 Higgins et al., 2017] and generally show weaker generative ability than GANs [Goodfellow et al.,  
 42 2020] or diffusion models [e.g., DDPM Ho et al., 2020].

43 Researchers have explored combining VAEs with diffusion models to improve generative quality, for  
 44 example by learning a diffusion prior [Wehenkel and Louppe, 2021] or training diffusion models on  
 45 VAE latent spaces [Kwon et al., 2022, Yan et al., 2021]. Beyond VAEs, masked autoencoders [MAEs,  
 46 He et al., 2022] have recently gained attention as efficient learners for images and videos. MAEs  
 47 reconstruct masked regions from the unmasked context, a strategy well-suited to modalities with  
 48 strong local structure like images or audio, but less effective for data with long-range dependencies  
 49 (researchers have also suggested that diffusion models can be interpreted as MAEs; Wei et al. 2023).  
 50 Despite their impressive performance in image and audio domains, these methods often struggle to  
 51 encode high-frequency structure in irregularly sampled long sequences, such as those obtained by  
 52 astronomical spectrographs.

### 53 3 Diffusion autoencoder with perceiver

54 In this section, we introduce our **diffusion autoencoder with perceiver (daep<sup>1</sup>)**. A unimodal daep has  
 55 three components: tokenizer, encoder, and diffusion decoder.

56 **Tokenizers.** We represent raw data as a sequence of tokens in the model dimension. We treat data  
 57 as a collection of measurements at specific locations with accompanying metadata. Formally, we  
 58 define  $(v, s, m)$ , where  $v$  are measurement values (e.g., flux of an astrophysical source),  $s$  is location  
 59 information (e.g., wavelength, time, band), and  $m$  is observational metadata (e.g., instrument,  
 60 observation time of spectra). We adapt the perceiver strategy [Jaegle et al., 2021b] by linearly  
 61 projecting  $v$  using fixed sinusoidal embeddings for continuous parts of  $s$  (e.g., time) and categorical  
 62 embeddings for discrete parts (e.g., color bands). We concatenate value and positional embeddings  
 63 and project them to model dimension. We represent metadata as extra tokens appended to the  
 64 sequence. For images, we add a small CNN at the beginning of the encoder to reduce the number of  
 65 tokens. We show three example tokenizers in fig. 1.

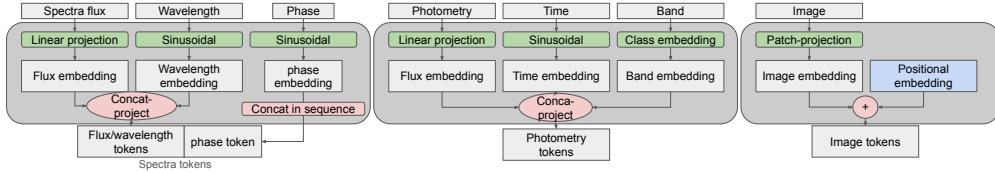


Figure 1: Tokenizers used in our empirical studies, from left to right: spectra (flux across wavelengths), light curves (brightness in different colors over time), and images.

66 **Unimodal encoders.** We use perceiver encoders [Jaegle et al., 2021b] to map token sequences into  
 67 compact bottleneck representations. Tokens act as Keys and Values in cross attention, while bottleneck  
 68 representations serve as Queries. Self-attention is applied only among bottleneck sequences. We  
 69 repeat these perceiver blocks several times, optionally sharing weights. This design handles variable-  
 70 length sequences with linear cost in sequence length, making it efficient for processing long and  
 71 irregular data. Finally, we project bottleneck sequences from the model dimension to a fixed  
 72 bottleneck dimension. We illustrate the encoder in fig. 2.

73 **Perceiver IO-based diffusion decoder.** Our decoder builds on diffusion transformers [Peebles and  
 74 Xie, 2023], particularly cross-attention conditioning. We encode diffusion time with fixed sinusoidal  
 75 embeddings passed through an MLP, as in Peebles and Xie [2023], and concatenate it with the

<sup>1</sup>Code available here.

76 conditioning representation for the score model. The score model, which predicts added noise, is a  
 77 perceiver-IO: noisy data is tokenized, concatenated with conditioning tokens, and used as Keys and  
 78 Values in cross attention. A latent sequence serves as Queries with self-attention, then acts as Keys  
 79 and Values in a second cross-attention stage with positional information as Queries. We repeat these  
 80 blocks, optionally sharing weights. The schematic is shown in fig. 2. While Jaegle et al. [2021a]  
 81 recommend latent lengths of 128–512, this can be longer than data sequence for some tasks. In such  
 82 cases, we use a single-stage perceiver decoder without a latent sequence, directly connecting noisy  
 83 tokens to noise prediction through cross attention.

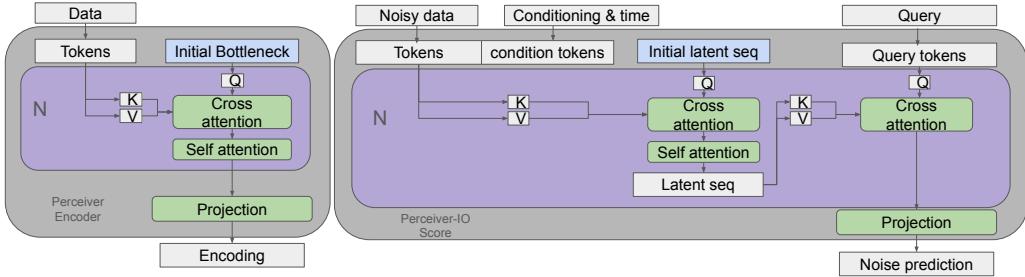


Figure 2: Schematic of the perceiver encoder and perceiver-IO score model used in daep.

84 **Training and sampling.** We train with the score-matching loss from DDPM [Ho et al., 2020] using  
 85 1,000 denoising steps. At inference, we adopt deterministic DDIM [Song et al., 2020] for faster  
 86 sampling with 200 steps. Similar to Preechakul et al. [2022], our model is not a generative model  
 87 since it requires the bottleneck representation. However, following Preechakul et al. [2022] and  
 88 Wehenkel and Louppe [2021], we can train another DDIM to sample from the bottleneck distribution,  
 89 enabling prior generation.

## 90 4 Unimodal experiments

91 **High-resolution spectra of variable stars.** We used data from v2.0 DR9 of the Large Sky Area  
 92 Multi-Object Fiber Spectroscopic Telescope [LAMOST, Cui et al., 2012], specifically the dataset  
 93 consolidated by Rizhko and Bloom [2025]. The dataset contains spectra of variable stars with on  
 94 average  $\sim$ 2,500 measurements and up to  $\sim$ 4,000 per star. We trained our model and a benchmark  
 95  $\beta$ -VAE ( $\beta = 0.1$ ) with the same perceiver encoder and decoder to encode spectra into a four-token  
 96 sequence of dimension eight. Full architectural details are provided in appendix A.1. We show  
 97 two enlarged test examples in fig. 3, with additional examples in fig. 6. Our model reconstructions  
 98 captured finer spectral features compared to the VAE baseline.

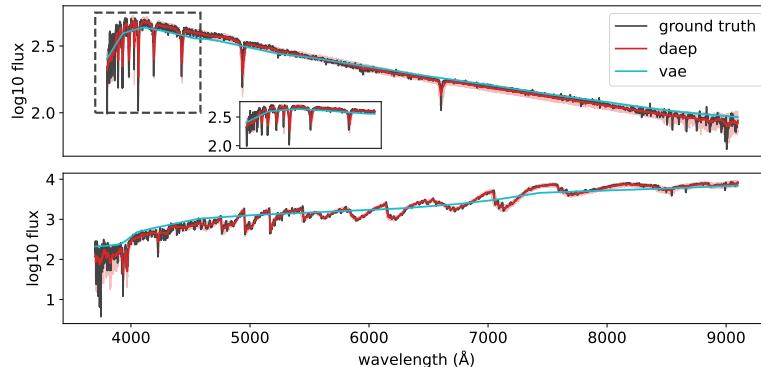


Figure 3: Two example reconstructions of long variable star spectra. Our daep model captured finer spectroscopic features, while the VAE mostly reproduced the overall continuum, likely due to posterior collapse.

99 **Spectra and photometry of supernovae.** We used data from the Zwicky Transient Facility Bright  
100 Transient Survey [ZTFBTS, Bellm et al., 2018]. In this survey, supernovae are measured in two  
101 bands—green (g) and red (r)—along with spectra, all sampled irregularly. We encoded the light  
102 curves into a two-token sequence of dimension two. Example reconstructions from our model are  
103 shown in fig. 4, with VAE results in fig. 7. Our method achieved more accurate reconstructions than  
104 the VAE baseline. Implementation details are provided in appendix A.2, with additional spectral  
105 results in appendix B.3.

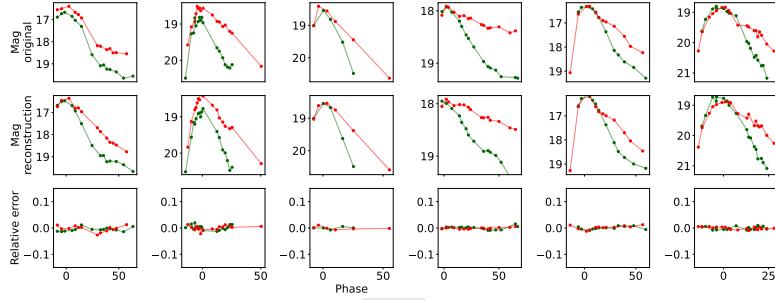


Figure 4: Light curve reconstruction from two latent tokens of dimension two using our daep model.

106 **Images of galaxies.** We used data from the Galaxy10 DECals dataset. Our method achieves  
107 moderately superior reconstructions compared to the VAE baseline, particularly for high-frequency  
108 features. Further details and results are provided in appendix A.3 and appendix B.4.

## 109 5 Towards multimodality

110 **Modality mixing and training for multimodal data.** To learn joint representations from multiple  
111 modalities, we used a late mixing strategy. We first encoded each modality with a perceiver encoder,  
112 added a learnable modality embedding, concatenated them along the sequence dimension, and then  
113 applied another perceiver encoder as a “mixer” to produce a single compact bottleneck sequence.  
114 Because the perceiver encoder does not require fixed-length input, we trained with modality dropping  
115 [Neverova et al., 2015, Liu et al., 2022] so the multimodal model can handle missing modalities.  
116 We always decoded all modalities using modality-specific diffusion decoders. A schematic of this  
117 architecture with two modalities is shown in fig. 5.

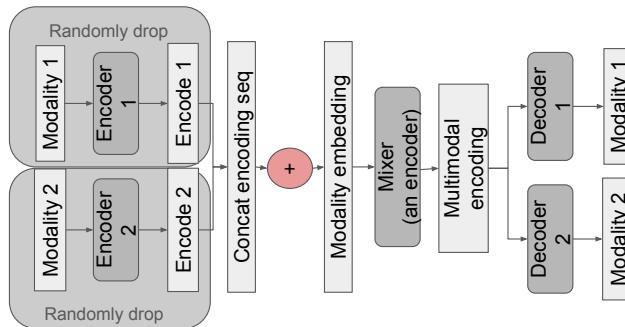


Figure 5: Late mixing and modality dropping for the multimodal daep model.

118 **Simulated supernova spectra and photometry.** We used data from simulated, idealized type Ia  
119 supernova light curves and spectra provided by Goldstein and Kasen [2018] and Shen and Gagliano  
120 [2025]. We trained a multimodal daep and evaluated it on cross-modality generation. Details are  
121 deferred to appendix A.4 and appendix B.5.

122 **References**

- 123 Behzad Abareshi, J Aguilar, S Ahlen, Shadab Alam, David M Alexander, R Alfarsy, L Allen,  
124 C Allende Prieto, O Alves, J Ameel, et al. Overview of the instrumentation for the dark energy  
125 spectroscopic instrument. *The Astronomical Journal*, 164(5):207, 2022.
- 126 P\_D Aleo, K Malanchev, S Sharief, D\_O Jones, G Narayan, R\_J Foley, V\_A Villar, C\_R Angus, V\_F  
127 Baldassare, M\_J Bustamante-Rosell, et al. The young supernova experiment data release 1 (yse  
128 dr1): light curves and photometric classification of 1975 supernovae. *The Astrophysical Journal  
129 Supplement Series*, 266(1):9, 2023.
- 130 Eric C Bellm, Shrinivas R Kulkarni, Matthew J Graham, Richard Dekany, Roger M Smith, Reed  
131 Riddle, Frank J Masci, George Helou, Thomas A Prince, Scott M Adams, et al. The zwicky  
132 transient facility: system overview, performance, and first results. *Publications of the Astronomical  
133 Society of the Pacific*, 131(995):018002, 2018.
- 134 Xiang-Qun Cui, Yong-Heng Zhao, Yao-Quan Chu, Guo-Ping Li, Qi Li, Li-Ping Zhang, Hong-Jun  
135 Su, Zheng-Qiu Yao, Ya-Nan Wang, Xiao-Zheng Xing, et al. The large sky area multi-object fiber  
136 spectroscopic telescope (lamost). *Research in Astronomy and Astrophysics*, 12(9):1197, 2012.
- 137 Richard Dekany, Roger M Smith, Reed Riddle, Michael Feeney, Michael Porter, David Hale, Jeffry  
138 Zolkower, Justin Belicki, Stephen Kaye, John Henning, et al. The zwicky transient facility:  
139 Observing system. *Publications of the Astronomical Society of the Pacific*, 132(1009):038001,  
140 2020.
- 141 Prafulla Dhariwal and Alexander Nichol. Diffusion models beat gans on image synthesis. *Advances  
142 in neural information processing systems*, 34:8780–8794, 2021.
- 143 Daniel A Goldstein and Daniel Kasen. Evidence for sub-chandrasekhar mass type ia supernovae  
144 from an extensive survey of radiative transfer models. *The Astrophysical Journal Letters*, 852(2):  
145 L33, 2018.
- 146 Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair,  
147 Aaron Courville, and Yoshua Bengio. Generative adversarial networks. *Communications of the  
148 ACM*, 63(11):139–144, 2020.
- 149 Matthew J Graham, SR Kulkarni, Eric C Bellm, Scott M Adams, Cristina Barbarino, Nadejda  
150 Blagorodnova, Dennis Bodewits, Bryce Bolin, Patrick R Brady, S Bradley Cenko, et al. The  
151 zwicky transient facility: science objectives. *Publications of the Astronomical Society of the Pacific*,  
152 131(1001):078001, 2019.
- 153 Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked  
154 autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF conference on computer  
155 vision and pattern recognition*, pages 16000–16009, 2022.
- 156 Irina Higgins, Loic Matthey, Arka Pal, Christopher Burgess, Xavier Glorot, Matthew Botvinick,  
157 Shakir Mohamed, and Alexander Lerchner. beta-vae: Learning basic visual concepts with a  
158 constrained variational framework. In *International conference on learning representations*, 2017.
- 159 Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in  
160 neural information processing systems*, 33:6840–6851, 2020.
- 161 Andrew Jaegle, Sebastian Borgeaud, Jean-Baptiste Alayrac, Carl Doersch, Catalin Ionescu, David  
162 Ding, Skanda Koppula, Daniel Zoran, Andrew Brock, Evan Shelhamer, et al. Perceiver io: A  
163 general architecture for structured inputs & outputs. *arXiv preprint arXiv:2107.14795*, 2021a.
- 164 Andrew Jaegle, Felix Gimeno, Andy Brock, Oriol Vinyals, Andrew Zisserman, and Joao Carreira.  
165 Perceiver: General perception with iterative attention. In *International conference on machine  
166 learning*, pages 4651–4664. PMLR, 2021b.
- 167 DO Jones, RJ Foley, G Narayan, Jens Hjorth, ME Huber, PD Aleo, KD Alexander, CR Angus, Katie  
168 Auchettl, VF Baldassare, et al. The young supernova experiment: survey goals, overview, and  
169 operations. *The Astrophysical Journal*, 908(2):143, 2021.

- 170 Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- 172 Mingi Kwon, Jaeseok Jeong, and Youngjung Uh. Diffusion models already have a semantic latent space. *arXiv preprint arXiv:2210.10960*, 2022.
- 174 Han Liu, Yubo Fan, Hao Li, Jiacheng Wang, Dewei Hu, Can Cui, Ho Hin Lee, Huahong Zhang, and Ipek Oguz. Moddrop++: A dynamic filter network with intra-subject co-training for multiple sclerosis lesion segmentation with missing modalities. In *International Conference on Medical Image Computing and Computer-Assisted Intervention*, pages 444–453. Springer, 2022.
- 178 Frank J Masci, Russ R Laher, Ben Rusholme, David L Shupe, Steven Groom, Jason Surace, Edward Jackson, Serge Monkewitz, Ron Beck, David Flynn, et al. The zwicky transient facility: Data processing, products, and archive. *Publications of the Astronomical Society of the Pacific*, 131(995):018003, 2018.
- 182 Natalia Neverova, Christian Wolf, Graham Taylor, and Florian Nebout. Moddrop: adaptive multi-modal gesture recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(8):1692–1706, 2015.
- 185 William Peebles and Saining Xie. Scalable diffusion models with transformers. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 4195–4205, 2023.
- 187 Konpat Preechakul, Nattanat Chatthee, Suttisak Wizadwongsu, and Supasorn Suwajanakorn. Diffusion autoencoders: Toward a meaningful and decodable representation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10619–10629, 2022.
- 190 Ali Razavi, Aaron Van den Oord, and Oriol Vinyals. Generating diverse high-fidelity images with vq-vae-2. *Advances in neural information processing systems*, 32, 2019.
- 192 Maria Rizhko and Joshua S Bloom. Astrom3: A self-supervised multimodal model for astronomy. *The Astronomical Journal*, 170(1):28, 2025.
- 194 Yunyi Shen and Alexander T Gagliano. Variational diffusion transformers for conditional sampling of supernovae spectra. *arXiv preprint arXiv:2505.03063*, 2025.
- 196 Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. *arXiv preprint arXiv:2010.02502*, 2020.
- 198 J. L. Tonry, L. Denneau, A. N. Heinze, B. Stalder, K. W. Smith, S. J. Smartt, C. W. Stubbs, H. J. Weiland, and A. Rest. ATLAS: A High-cadence All-sky Survey System. *Publications of the Astronomical Society of the Pacific*, 130(988):064505, June 2018. doi: 10.1088/1538-3873/aabadf.
- 201 Arash Vahdat and Jan Kautz. Nvae: A deep hierarchical variational autoencoder. *Advances in neural information processing systems*, 33:19667–19679, 2020.
- 203 Aaron Van Den Oord, Oriol Vinyals, et al. Neural discrete representation learning. *Advances in neural information processing systems*, 30, 2017.
- 205 Antoine Wehenkel and Gilles Louppe. Diffusion priors in variational autoencoders. *arXiv preprint arXiv:2106.15671*, 2021.
- 207 Chen Wei, Karttikeya Mangalam, Po-Yao Huang, Yanghao Li, Haoqi Fan, Hu Xu, Huiyu Wang, Cihang Xie, Alan Yuille, and Christoph Feichtenhofer. Diffusion models as masked autoencoders. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 16284–16294, 2023.
- 211 Wilson Yan, Yunzhi Zhang, Pieter Abbeel, and Aravind Srinivas. Videogpt: Video generation using vq-vae and transformers. *arXiv preprint arXiv:2104.10157*, 2021.
- 213 Donald G York, Jennifer Adelman, John E Anderson Jr, Scott F Anderson, James Annis, Neta A Bahcall, JA Bakken, Robert Barkhouser, Steven Bastian, Eileen Berman, et al. The Sloan digital sky survey: Technical summary. *The Astronomical Journal*, 120(3):1579, 2000.
- 216 Gemma Zhang, Thomas Helfer, Alexander T Gagliano, Siddharth Mishra-Sharma, and V Ashley Villar. Maven: a multimodal foundation model for supernova science. *Machine Learning: Science and Technology*, 5(4):045069, 2024.

219 **A Implementation details**

220 **A.1 LAMOST model details**

221 **Data preprocessing** We enforced a  $3\sigma$  quality cut, i.e., only measurements exceeding 3 times the  
 222 measurement error are kept for modeling. After all quality cuts, we have 17,063 training stars. We  
 223 took arcsinh of flux before modeling and after generation we calculate the original flux. Flux and  
 224 wavelength are standardized to a z-score using the mean and standard deviation of the whole training  
 225 set after calculating arcsinh of flux.

**Architectural details Training details** We used same learning rate of  $2.5e - 4$  for both models and

model	bottleneck len	bottleneck dim	enc. layers	dec. layers	model dim	# heads	hidden seq len
daep	4	8	4	4	128	8	256
VAE	4	8	4	4	128	8	256

Table 1: Architectural choices used in LAMOST experiments.

226 trained for 2000 epochs and 200 epochs respectively for daep and VAE, both training loss converged.  
 227 We set  $\beta = 0.1$  for VAE.

229 **A.2 ZTF model details**

230 **Light curve preprocessing** We first enforced a  $3\sigma$  cut on measurements, then used a Gaussian  
 231 process to find the peak time of red band as the 0 phase. We align time to be relative to the peak time.  
 232 We only kept events that light curves cover the peak.

233 **Spectra preprocessing** We enforced a  $3\sigma$  quality cut for both spectra and lightcurve, i.e., only  
 234 measurements exceeding 3 times the measurement error are kept for modeling. After all cuts we have  
 235 2,934 events left in training. We took the base-10 logarithm of the flux before modeling, and after  
 236 generation we calculate the original flux. We also apply a median filter to filter out noise. Flux and  
 237 wavelength are then standardized to a z-score using the mean and standard deviation of the whole  
 238 training set.

**Architectural details**

model	bottleneck len	bottleneck dim	enc. layers	dec. layers	model dim	# heads
daep	4	4	4	4	128	4
VAE	4	4	4	4	128	4

Table 2: Architectural choices used in ZTF spectra experiments. We used a single stage decoder  
 (skipping latent sequence) since the sequence is short.

239

model	bottleneck len	bottleneck dim	enc. layers	dec. layers	model dim	# heads
daep	2	2	4	4	128	4
VAE	2	2	4	4	128	4

Table 3: Architectural choices used in ZTF light curve experiments. We used a single stage decoder  
 (skipping latent sequence) since the sequence is short.

240 **Training details** Different from LAMOST experiment, we augment our data by 5 folds and adding  
 241 noise to flux measurement and randomly mask part of the measurements. We used same learning rate  
 242 of  $2.5e - 4$  for both models and trained for 2000 epochs and 200 epochs respectively for daep and  
 243 VAE, both training loss converged. We set  $\beta = 0.1$  for VAE.

244 **A.3 Galaxy10 model details**

245 **Preprocessing** We normalize the pixel values assuming mean 0.5 and std 0.5 and reshape to size  
 246  $64 \times 64 \times 3$ .

247 **Architectural details**

model	bottleneck len	bottleneck dim	enc. layers	dec. layers	model dim	# heads	patch size
daep	8	8	4	4	128	4	4
VAE	8	8	4	4	128	4	4

Table 4: Architectural choices used in Galaxy10 experiments. We used a single stage decoder (skipping latent sequence) since the sequence is short.

248 **Training details** We augment our data by 3 folds with random flipping. We used same learning rate  
 249 of  $2.5e - 4$  for both models and trained for 2000 epochs and 200 epochs respectively for daep and  
 250 VAE, both training loss converged. We set  $\beta = 0.5$  for VAE.

#### 251 A.4 Multimodal spectra and photometry

252 **Data preprocessing.** We did not perform further processing beyond those used in Shen and Gagliano  
 253 [2025].

254 **Architectural details** We have the first stage encoder for both light curve and spectra to have model  
 255 dimension 256, 4 layers, 4 heads, and encode to 64 tokens. The mixer has 4 layers and 4 heads and  
 256 model dimension 256, during encoder we allow the concatenated sequence to have self attention. We  
 257 encode to a bottleneck sequence of 4 tokens of dimension 4 each.

258 **Training details.** In each batch we randomly dropped a modality with probability 0.2, but retain at  
 259 least one modality. We trained with learning rate  $2.5e - 4$  and for 2000 epochs. The loss converged.

## 260 B Further experimental results

### 261 B.1 More high-resolution spectra

262 In fig. 6, we show additional spectra reconstructions using daep and VAE baselines. Our method  
 263 consistently captures higher-frequency information details compared to the VAE baseline.

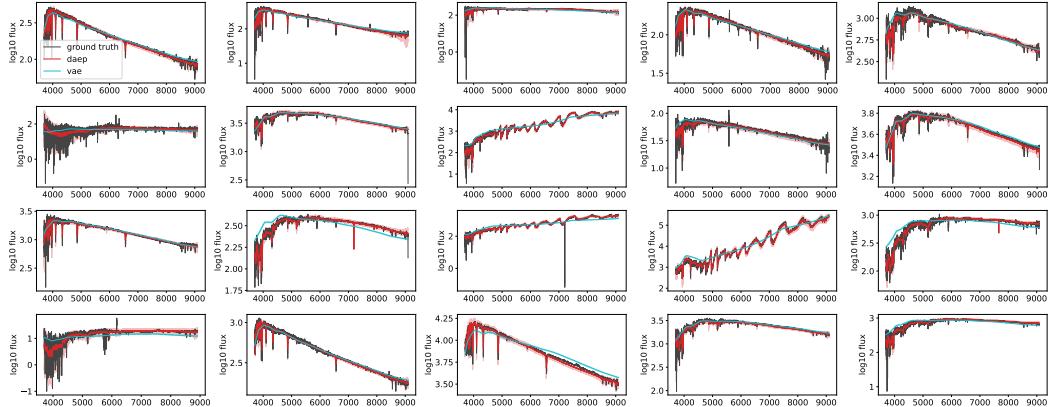


Figure 6: Reconstructions of additional LAMOST variable star spectra. Our method (red) captures more high-frequency absorption features than the VAE baseline with the same-sized bottleneck representation (blue).

### 264 B.2 ZTF light curves

265 In fig. 7, we show ZTF light curve reconstructions with the VAE baseline. Our method performs  
 266 moderately better.

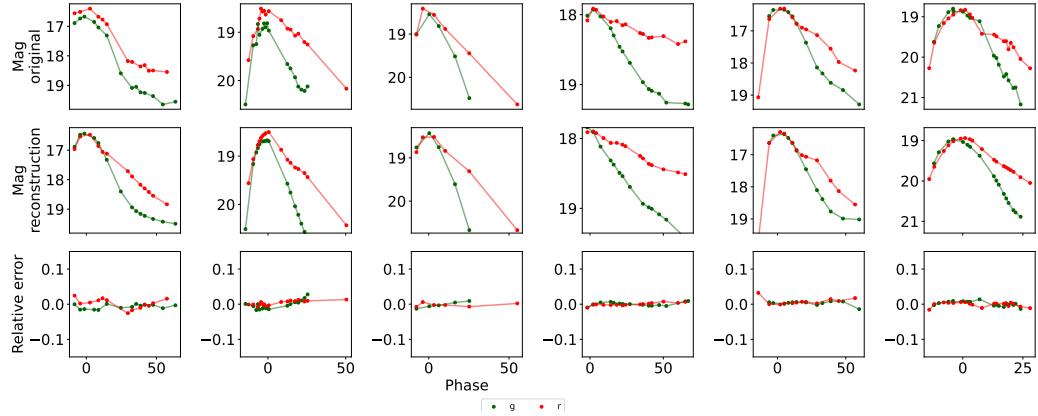


Figure 7: Reconstruction of ZTF light curves using a baseline VAE model (middle row) compared to the ground truth observations (top row). Residuals are given in the bottom row.

267    **B.3 ZTF spectra**

268    In figs. 8 and 9, we show spectra reconstructions on ZTF data. Our method captures finer details and  
269    produces better-covered posteriors than the VAE baseline.

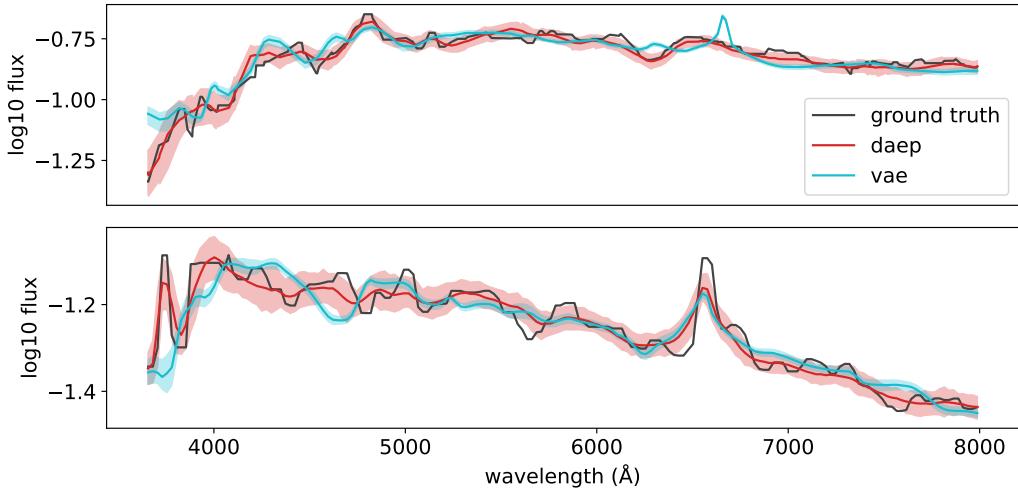


Figure 8: Zoomed-in ZTF spectra reconstructions with daep and VAE. Our method captures finer details and maintains better posterior coverage.

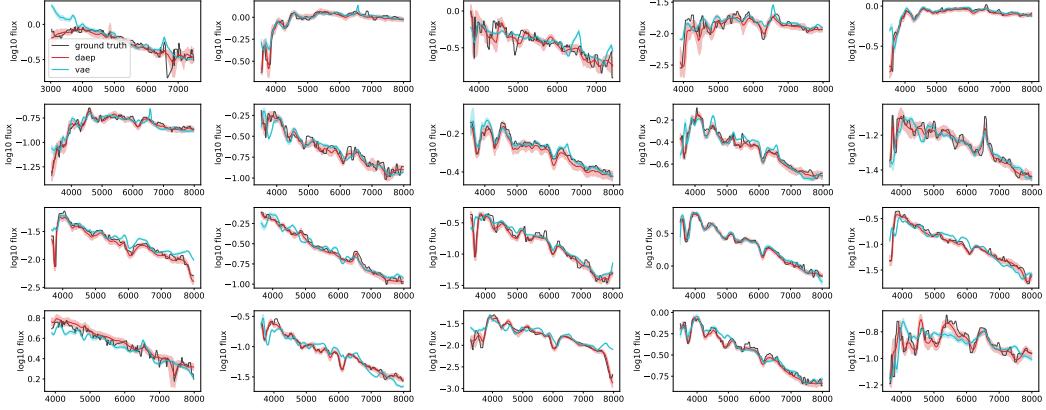


Figure 9: Additional ZTF spectra reconstructions with daep and VAE. Our method captures finer details and maintains better posterior coverage.

270 In fig. 10, we compare latent representations (after t-SNE) from daep and VAE, colored by event type.  
 271 Interestingly, the daep latent space appears more continuous, while the VAE with  $\beta = 0.1$  shows  
 272 more holes.

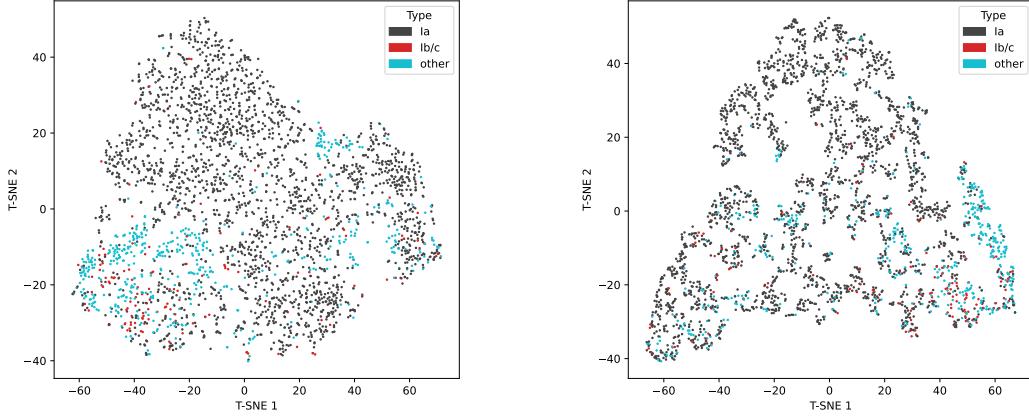


Figure 10: Latent representations of ZTF spectra from daep (left) and VAE (right). The latent space for daep appears more well-regularized compared to a VAE of comparable dimensionality.

#### 273 B.4 Galaxy10

274 In figs. 11 and 12, we show reconstructions on the Galaxy10 dataset. Our method captures slightly  
 275 finer-scale structures such as spiral arms compared to the VAE baseline.

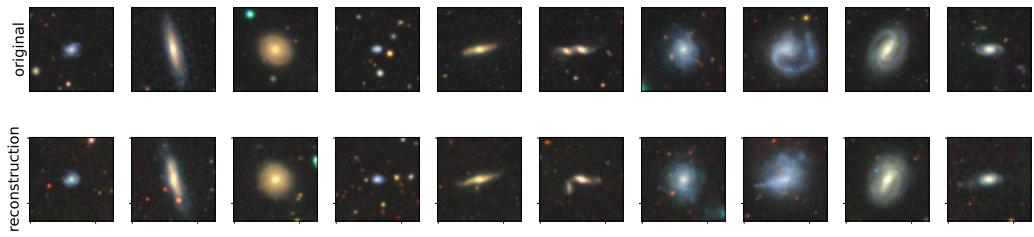


Figure 11: Reconstructions on Galaxy10 test data using daep.

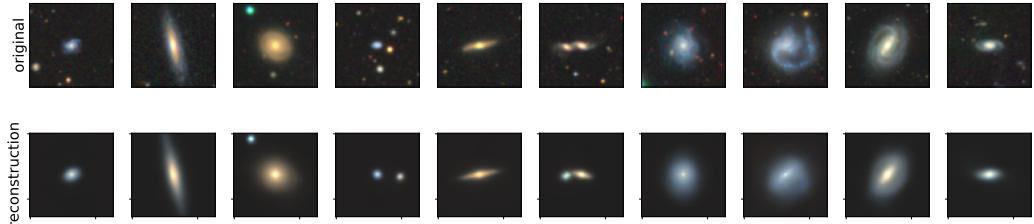


Figure 12: Reconstructions on Galaxy10 test data using a VAE. Reconstructions are heavily smoothed and retain fewer high-frequency structure than daep reconstructions.

276 **B.5 Multimodal spectra and LSST photometry of supernovae**

277 We demonstrate cross-modality generation using the learned encoder. In fig. 13, we show inference  
 278 from light curves to spectra, while in fig. 14, we show inference from spectra to light curves. Our  
 279 model performs cross-modality inference directly from the bottleneck representation.

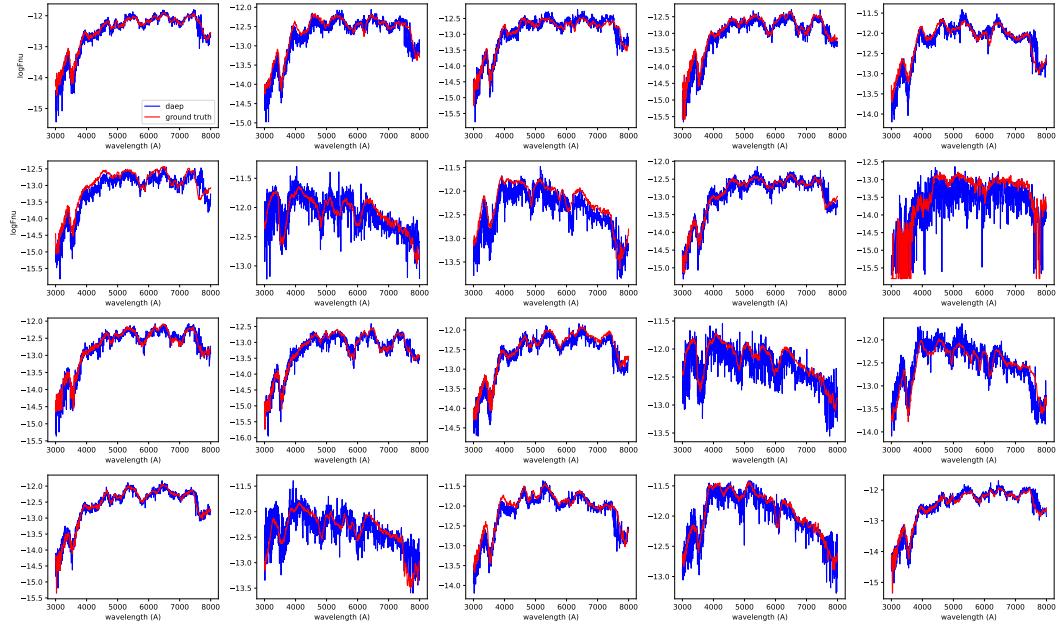


Figure 13: Cross-modality inference from light curves to spectra in the Goldstein simulation.

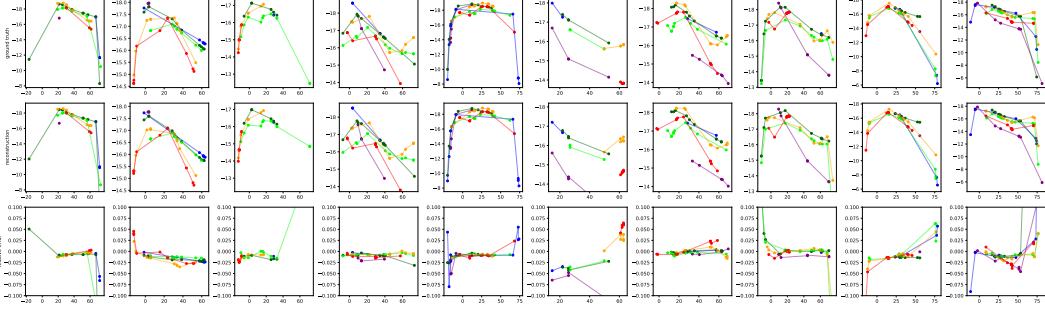


Figure 14: Cross-modality inference from spectra to light curves in the Goldstein simulation.

## 280 C Discussion

281 We have presented a perceiver-based diffusion autoencoder that scalably learns bottleneck representations  
 282 from irregularly sampled sequences. Its latent spaces appear more regular than those of a  
 283  $\beta$ -VAE baseline (Figure 10), while achieving better reconstruction on multiple time-domain datasets.  
 284 Further, the results for cross-modality generation show that the daep architecture can be valuable for  
 285 the development of multi-modal models in astrophysics. In the future, we plan to further evaluate the  
 286 learned representations on downstream tasks such as classification and clustering.