
Re-envisioning *Euclid* Galaxy Morphology: Identifying and Interpreting Features with Sparse Autoencoders

John F. Wu*

Space Telescope Science Institute
Johns Hopkins University
jfwu@stsci.edu

Michael Walmsley*

University of Toronto
m.walmsley@utoronto.ca

Abstract

Sparse Autoencoders (SAEs) can efficiently identify candidate monosemantic features from pretrained neural networks for galaxy morphology. We demonstrate this on Euclid Q1 images using both supervised (Zoobot) and new self-supervised (MAE) models. Our publicly released MAE achieves superhuman image reconstruction performance. While a Principal Component Analysis (PCA) on the supervised model primarily identifies features already aligned with the Galaxy Zoo decision tree, SAEs can identify interpretable features outside of this framework. SAE features also show stronger alignment than PCA with Galaxy Zoo labels. Although challenges in interpretability remain, SAEs provide a powerful engine for discovering astrophysical phenomena beyond the confines of human-defined classification.

1 Introduction

Galaxy morphology concepts are traditionally encoded in manually-designed taxonomies (e.g., De Vaucouleurs, 1963). These may miss concepts that are too rare to be manually detected. New, rare concepts are statistically guaranteed to be represented in the billions of resolved galaxies soon to be imaged by space telescopes like *Euclid* and *Roman* (Scaramella et al., 2021; Spergel et al., 2015), so we will need a data-driven method to identify these rare concepts. Relatedly, the scale of *Euclid* and *Roman* also demands an increasing reliance on deep learning models for interpreting images (e.g. classification, segmentation, feature extraction for multimodal tasks, etc.) Identifying the concepts learned by our models can inform how we build them and help mitigate the risk of unintended biases or shortcut learning (Geirhos et al., 2020).

Sparse autoencoders (SAEs) are routinely used for interpreting language models by identifying the set of directions in activation space that can together describe any activation vector in a neural network (Makhzani and Frey, 2014). Sparsity encourages these directions to be monosemantic, i.e., corresponding to a single concept (Elhage et al., 2022; Cunningham et al., 2023; Wu, 2025; Karvonen et al., 2025). Here, we apply SAEs to reveal the concepts learned by galaxy morphology models.

We construct embeddings for galaxies in the *Euclid* Q1 data release (Euclid Collaboration et al., 2025a) and then compress those embeddings via SAEs and PCA to identify learned concepts. We show that SAEs efficiently extract candidate morphological features from pretrained networks, comparing against a PCA baseline. For both supervised and self-supervised embeddings, SAEs identify interpretable features aligned with, and beyond the Galaxy Zoo decision trees, while PCA primarily aligns with existing categories.

*The authors contributed equally to this work.

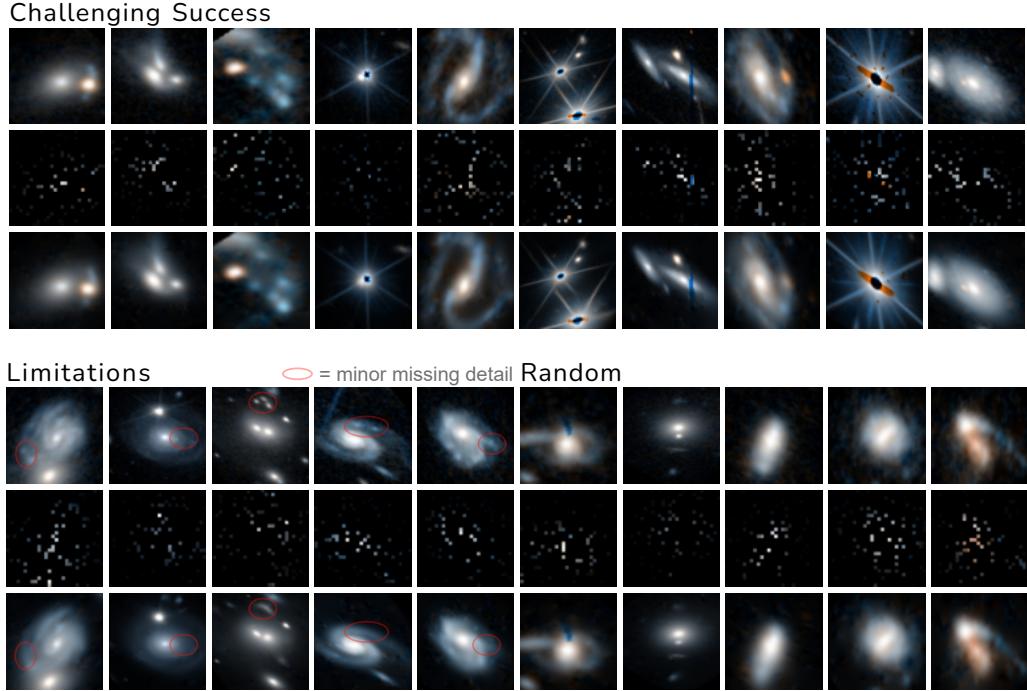


Figure 1: True images (*top row*), masked inputs (*middle row*), and reconstructed outputs (*bottom row*), grouped as follows. Top: challenging images with successful MAE reconstructions. Lower left: images with missing local details (e.g., small starforming clumps, multiple background sources with marginal separation). These cannot be reconstructed if not at least partially included in an unmasked patch. Masking is only applied during training and so these details are still in the embeddings. Lower right: random images (first five) demonstrating that near-perfect reconstructions are the norm.

2 Representations of Galaxy Morphology

Euclid Imaging and Embeddings. We analyze two types of neural network embeddings for galaxy images in the Euclid Q1 dataset. For supervised embeddings, we use the Zoobot models presented in [Euclid Collaboration et al. \(2025c\)](#). These models are pretrained on approximately 1M volunteer-annotated galaxies from four other telescopes and then finetuned using 170k Euclid galaxies with new annotations from Galaxy Zoo Euclid (GZ). Our chosen model, ConvNeXT-Nano, has an encoder with 640-dimensional features.

For self-supervised embeddings, we train a masked autoencoder (MAE, [He et al. 2022](#)) on 3M Euclid images drawn from the internal Euclid dataset “RR2”. This is $5 - 10 \times$ larger than recent work applying MAE to astronomy images ([Fathkouhi and Fox, 2024](#)) and applying other self-supervised learning approaches to Euclid images ([Euclid Collaboration et al., 2025b](#)).

Our MAE is a ViT-S¹ encoder (30.1M parameters, 384-dimensional features) plus a 3-layer decoder, trained for 2 A100-weeks. We make several design changes compared to the standard ConvNeXT/MAE approach in order to achieve stronger reconstruction performance:

We use a high masking fraction of 90% (vs. 75% as published). We find that networks of just 1M parameters perform well with conventional masking ratios, suggesting that astronomy images have high redundancy. Work using MAE for video (which is intrinsically high redundancy) shows that higher masking ratios are required to force the network to do more than interpolate ([Feichtenhofer et al., 2022](#)).

¹https://huggingface.co/timm/vit_small_patch16_224.augreg_in21k

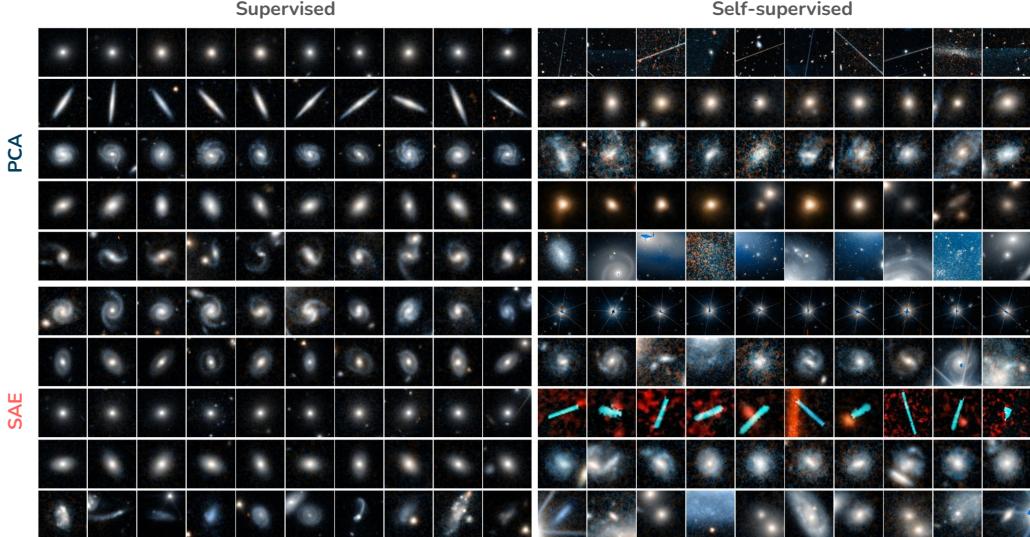


Figure 2: Top 10 examples for first five features extracted via PCA (*upper*), and SAEs (*lower*), for embeddings extracted from supervised (*left*) and self-supervised (*right*) models.

We use a small patch area of 8×8 pixels (vs. 16×16 as published). Astronomy images are relatively sparse, often with large ‘blank’ areas and compact background/foreground galaxies. Random sampling with large patches misses relevant areas entirely, making prediction impossible.

We adjust the frequency of the positional encoding to indicate the angular scale of each image. This is an astronomical analogy to work using MAE for remote sensing (Reed et al., 2023), who adjust the frequency to indicate the ground scale of each satellite image.

Our trained MAE reconstructs masked galaxy features more accurately than professional astronomers. Figure 1 shows our MAE infilling 90% masked galaxy images with near-perfect accuracy. An informal survey of our human colleagues suggests that professional astronomers cannot correctly describe the same masked images. *Superhuman* reconstruction performance indicates that our MAE extracts meaningful features. We also hope that our MAE will be useful for tasks like patch-level anomaly search (find me images that include patches like *these*) and targeted infill (remove *this* small artifact). We release our MAE model on HuggingFace at <https://huggingface.co/mwalmley/euclid-rr2-mae>. We also present a demo of the MAE at https://huggingface.co/spaces/mwalmley/euclid_masked_autoencoder.

Extracting Features from Embeddings. SAEs learn sparse representations of the embeddings using an *overcomplete* basis. We train a Matryoshka Sparse Autoencoder (SAE; (Bussmann et al., 2025)) with hierarchical group sizes $64, 64, 128, \dots, 1024$, batch top- k sparsity ($k = 64$), a small L1 penalty, and auxiliary loss for reviving dead neurons (following (O’Neill et al., 2024)). For both self-supervised and supervised embeddings, we train for 200 epochs. The code is publicly available at <https://github.com/jwuphysics/euclid-galaxy-morphology-saes>.

We use top $k = 64$ components in terms of activation frequency. Figure 2 shows the first 5 for each. They explain $\sim 83\%$ of the variance for Zoobot and $\sim 90\%$ of the variance for the MAE. We also perform PCA decomposition of Zoobot and MAE embeddings; the cumulative explained variance using first 64 components is $\sim 90\%$ for Zoobot and $\sim 95\%$ for the MAE.

Interpreting SAE Features via Alignment with Galaxy Zoo. For each learned feature, we compute Spearman rank correlations (r) with all Galaxy Zoo morphology labels and take the maximum absolute correlation.² Higher values indicate stronger correlation with known morphological classifications. We report the mean over the top k PCA or SAE features, where “top” is defined as the highest rank-ordered eigenvalues for PCA, and the most frequently activated neurons for SAEs.

²See Appendix B for a similar analysis on whether SAE features are predictable from GZ vote fractions.

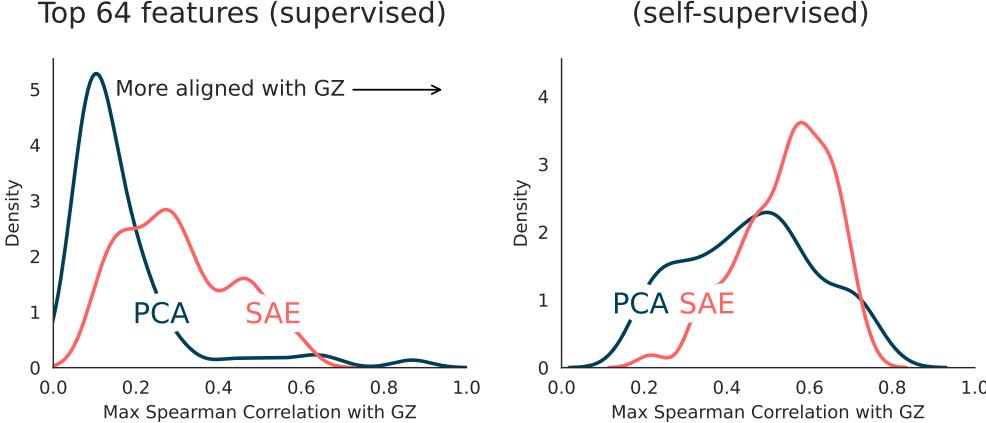


Figure 3: Top SAE features (red) are generally more aligned with GZ morphological classifications than PCA (blue) for both the supervised model embeddings (left) and self-supervised model embeddings (right). Both panels show the distribution of maximum Spearman rank correlation r values between top 64 features and any GZ class.

3 Results

In Figure 3, we show how well the SAE and PCA features correlate with GZ features. An extension of this analysis beyond the top $k = 64$ features is presented in the Appendix A (see, e.g., Figure 4).

Supervised: The top $k = 5$ principal components of the supervised embeddings strongly correlate with GZ features, even more so than the corresponding top 5 SAE features (see Table 1). When we extend to the top $k = 64$, higher-order PCA features show weaker alignment with GZ classes than the SAE features. At least, this implies that SAE features retain coherency while PCA features become more noisy at $k = 64$. We see SAE features outside of the GZ decision tree, e.g., dust lanes in edge-on disk galaxies, elliptical galaxies with bluer companions, etc.

All of the strongest correlations between PCA and SAE features are due to the smooth-or-featured-euclid_* node in the GZ decision tree. This is unsurprising, as the root node classification routes downstream morphological classifications into “secondary” categories. Thus we repeat our alignment tests by omitting this root node, and only using secondary. Overall, we find that alignment with GZ decreases, but we observe the same trends as before.

Self-supervised: The MAE embeddings exhibit surprisingly strong correlations with GZ—they are more correlated than the Zoobot embeddings. We suspect this arises because the MAE embeddings are intent on reconstructing imaging artifacts such as saturated stars or ghosts, which can be diverse and span many pixels. The SAE and PCA features primarily activate on these artifacts, and indeed we see that the most of the strongly aligned GZ classes are problem_* or artifact_*.³

4 Discussion

Limitations. Unfortunately, some SAE features remain difficult to interpret, despite the benefits of using Matryoshka SAEs to combat feature splitting and absorption (Chanin et al., 2024; Bussmann et al., 2025). Still, SAEs surpass our PCA baseline approach. PCA features become unidentifiable after ~ 30 features and show no alignment with GZ. In contrast, our SAE delivers an order of magnitude more aligned features (see Appendix A for more details).

³Another possibility is that the supervised optimization objective (i.e. aligning to GZ) can (counterintuitively) cram class-separating features into a low-dimensional subspace (e.g., Grigg et al., 2021). In other words, the supervised objective can cause embeddings to span a small number of known, dominant morphological classes, whereas the self-supervised case allows representations to disentangle and occupy a broader space of potentially novel features.

Table 1: Alignment between SAE/PCA Features and Galaxy Zoo vote fractions. We report alignment with “all” GZ classes, as well as the alignment with GZ classes beyond the most common classes (“secondary”), for the top $k = 5$ and $k = 64$ features extracted via PCA and SAEs. Random feature statistics are described in Appendix C.

Features	Supervised		Self-supervised	
	(all)	(secondary)	(all)	(secondary)
PCA ($k = 5$)	0.618 ± 0.149	0.582 ± 0.097	0.455 ± 0.199	0.444 ± 0.193
SAE ($k = 5$)	0.402 ± 0.154	0.371 ± 0.161	0.589 ± 0.067	0.589 ± 0.067
PCA ($k = 64$)	0.176 ± 0.138	0.166 ± 0.137	0.434 ± 0.174	0.396 ± 0.178
SAE ($k = 64$)	0.296 ± 0.129	0.283 ± 0.126	0.523 ± 0.123	0.509 ± 0.128
Random	0.041 ± 0.002	0.040 ± 0.002	0.0072 ± 0.0003	0.0071 ± 0.0003

Scientific Discovery. SAEs can extract interpretable features from both supervised and self-supervised models at scale (Gao et al., 2024). Sparse activations are steerable and can be readily interpreted through automated means (Cunningham et al., 2023; O’Neill et al., 2024; Ye et al., 2024). Although there are concerns about whether SAEs are effective at learning *known* concepts, they excel at identifying *new* features compared to other methods (Muhamed et al., 2024; Movva et al., 2025; Marks et al., 2025; Peng et al., 2025). Thus, SAEs have tremendous potential to aid scientific discovery by categorizing and sifting through next-generation astronomical datasets (see also, e.g., Lochner and Bassett 2021; O’Ryan and Gómez 2025). By extracting features from *self-supervised* embedding models, SAEs can act as scalable discovery engines that surface rare or anomalous astrophysical phenomena for future targeted investigation.

References

- Bart Bussmann, Noa Nabeshima, Adam Karvonen, and Neel Nanda. Learning multi-level features with matryoshka sparse autoencoders, 2025. URL <https://arxiv.org/abs/2503.17547>.
- David Chanin, James Wilken-Smith, Tom’avs Dulka, Hardik Bhatnagar, and Joseph Bloom. A is for absorption: Studying feature splitting and absorption in sparse autoencoders. *ArXiv*, abs/2409.14507, 2024. URL <https://api.semanticscholar.org/CorpusID:272827216>.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse Autoencoders Find Highly Interpretable Features in Language Models, October 2023. URL <http://arxiv.org/abs/2309.08600>. arXiv:2309.08600 [cs].
- Gerard De Vaucouleurs. Revised Classification of 1500 Bright Galaxies. *The Astrophysical Journal Supplement Series*, 8:31, 1963.
- Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah. Toy models of superposition, 2022. URL <https://arxiv.org/abs/2209.10652>.
- Euclid Collaboration, H. Aussel, I. Tereno, M. Schirmer, G. Alguero, B. Altieri, E. Balbinot, T. de Boer, P. Casenove, P. Corcho-Caballero, H. Furusawa, J. Furusawa, M. J. Hudson, K. Jahnke, G. Libet, J. Macias-Perez, N. Masoumzadeh, J. J. Mohr, J. Odier, D. Scott, T. Vassallo, G. Verdoes Kleijn, A. Zachei, N. Aghanim, A. Amara, S. Andreon, N. Auricchio, S. Awan, R. Azzollini, C. Baccigalupi, M. Baldi, A. Balestra, S. Bardelli, A. Basset, P. Battaglia, A. N. Belikov, R. Bender, A. Biviano, A. Bonchi, D. Bonino, E. Branchini, M. Brescia, J. Brinchmann, S. Camera, G. Cañas-Herrera, V. Capobianco, C. Carbone, V. F. Cardone, J. Carretero, S. Casas, F. J. Castander, M. Castellano, G. Castignani, S. Cavaudi, K. C. Chambers, A. Cimatti, C. Colodro-Conde, G. Congedo, C. J. Conselice, L. Conversi, Y. Copin, F. Courbin, H. M. Courtois, M. Cropper, J.-G. Cuby, A. Da Silva, R. da Silva, H. Degaudenzi, J. T. A. de Jong, G. De Lucia, A. M. Di Giorgio, J. Dinis, C. Dolding, H. Dole, M. Douspis, F. Dubath, C. A. J. Duncan, X. Dupac, S. Dusini, A. Ealet, S. Escoffier, M. Fabricius, M. Farina, R. Farinelli, F. Faustini, S. Ferriol, S. Fotopoulou, N. Fourmanoit, M. Frailis, E. Franceschi, P. Franzetti, S. Galeotta, K. George, W. Gillard, B. Gillis, C. Giocoli, P. Gómez-Alvarez, J. Gracia-Carpio, B. R. Granett, A. Grazian, F. Grupp, L. Guzzo, S. Gwyn, S. V. H. Haugan, O. Herent, J. Hoar, H. Hoekstra, M. S. Holliman, W. Holmes, I. M.

Hook, F. Hormuth, A. Hornstrup, P. Hudelot, S. Ilić, M. Jhabvala, B. Joachimi, E. Keihänen, S. Kermiche, A. Kiessling, B. Kubik, K. Kuijken, M. Kümmel, M. Kunz, H. Kurki-Suonio, O. Lahav, Q. Le Boule, A. M. C. Le Brun, D. Le Mignant, P. Liebing, S. Ligori, P. B. Lilje, V. Lindholm, I. Lloro, G. Mainetti, D. Maino, E. Maiorano, O. Mansutti, S. Marcin, O. Marggraf, K. Markovic, M. Martinelli, N. Martinet, F. Marulli, R. Massey, S. Maurogordato, H. J. McCracken, E. Medinaceli, S. Mei, M. Melchior, Y. Mellier, M. Meneghetti, E. Merlin, G. Meylan, A. Mora, M. Moresco, P. W. Morris, L. Moscardini, S. Mourre, R. Nakajima, C. Neissner, R. C. Nichol, S.-M. Niemi, J. W. Nightingale, T. Nutma, C. Padilla, S. Paltani, F. Pasian, J. A. Peacock, K. Pedersen, W. J. Percival, V. Pettorino, S. Pires, G. Polenta, J. E. Pollack, M. Poncet, L. A. Popa, L. Pozzetti, G. D. Racca, F. Raison, R. Rebolo, A. Renzi, J. Rhodes, G. Riccio, H.-W. Rix, E. Romelli, M. Roncarelli, E. Rossetti, B. Rusholme, R. Saglia, Z. Sakr, A. G. Sánchez, D. Sapone, B. Sartoris, M. Sauvage, J. A. Schewtschenko, P. Schneider, M. Scodellaro, A. Secroun, E. Sefusatti, G. Seidel, M. Seiffert, S. Serrano, P. Simon, C. Sirignano, G. Sirri, J. Skottfelt, A. Spurio Mancini, L. Stanco, J. Steinwagner, C. Surace, P. Tallada-Crespi, D. Tavagnacco, A. N. Taylor, H. I. Teplitz, N. Tessore, S. Toft, R. Toledo-Moreo, F. Torradeflot, A. Tsyganov, I. Tutusaus, E. A. Valentijn, L. Valenziano, J. Valiviita, A. Veropalumbo, Y. Wang, J. Weller, O. R. Williams, G. Zamorani, F. M. Zerbi, E. Zucca, V. Allevato, M. Ballardini, R. P. Blake, M. Bolzonella, E. Bozzo, C. Burigana, R. Cabanac, M. Calabrese, A. Cappi, D. Di Ferdinando, J. A. Escartin Vigo, L. Gabarra, W. G. Hartley, M. Huertas-Company, J. Martín-Fleitas, S. Matthew, M. Maturi, N. Mauri, R. B. Metcalf, A. Pezzotta, M. Pöntinen, C. Porciani, I. Risso, V. Scottez, M. Sereno, M. Tenti, M. Viel, M. Wiesmann, Y. Akrami, S. Alvi, I. T. Andika, S. Anselmi, M. Archidiacono, F. Atrio-Barandela, S. Avila, P. Bergamini, D. Bertacca, M. Bethermin, L. Bisigello, A. Blanchard, L. Blot, H. Böhringer, S. Borgani, A. S. Borlaff, M. L. Brown, S. Bruton, F. Buitrago, A. Calabro, G. Calderone, B. Camacho Quevedo, F. Caro, C. S. Carvalho, T. Castro, Y. Charles, F. Cogato, S. Conseil, A. R. Cooray, M. Costanzi, O. Cucciati, S. Davini, F. De Paolis, G. Desprez, A. Díaz-Sánchez, J. J. Diaz, S. Di Domizio, J. M. Diego, P. Dimauro, P.-A. Duc, A. Enia, Y. Fang, A. M. N. Ferguson, A. G. Ferrari, A. Finoguenov, A. Fontana, F. Fontanot, A. Franco, J. García-Bellido, T. Gasparotto, R. Gavazzi, E. Gaztanaga, F. Giacomini, F. Gianotti, A. H. Gonzalez, G. Gozaliasl, A. Gruppuso, M. Guidi, C. M. Gutierrez, A. Hall, C. Hernández-Monteagudo, H. Hildebrandt, J. Hjorth, J. Jacobson, S. Joudaki, J. J. E. Kajava, Y. Kang, V. Kansal, D. Karagiannis, K. Kiiveri, C. C. Kirkpatrick, S. Kruk, F. Lacasa, C. Laigle, M. Lattanzi, V. Le Brun, J. Le Graet, L. Legrand, M. Lembo, F. Lepori, G. Leroy, G. F. Lesci, J. Lesgourgues, L. Leuzzi, T. I. Liaudat, A. Loureiro, M. Magliocchetti, E. A. Magnier, C. Mancini, F. Mannucci, R. Maoli, C. J. A. P. Martins, L. Maurin, C. J. R. McPartland, J.-B. Melin, M. Migliaccio, M. Miluzio, P. Monaco, A. Montoro, C. Moretti, G. Morgante, C. Murray, S. Nadathur, K. Naidoo, A. Navarro-Alsina, S. Nesseris, L. Nicastro, M. Oguri, F. Paschalacqua, K. Paterson, L. Patrizii, A. Pisani, D. Potter, S. Quai, M. Radovich, P. Reimberg, P.-F. Rocci, G. Rodighiero, R. P. Rollins, S. Sacquegna, M. Sahlén, D. B. Sanders, E. Sarpa, C. Scarlata, J. Schaye, A. Schneider, M. Schultheis, D. Sciotti, D. Scognamiglio, E. Sellentin, F. Shankar, L. C. Smith, E. Soubrie, S. A. Stanford, K. Tanidis, C. Tao, G. Testera, M. Tewes, R. Teyssier, S. Tosi, A. Troja, M. Tucci, C. Valieri, A. Venhola, D. Vergani, F. Vernizzi, G. Verza, P. Vielzeuf, N. A. Walton, J. R. Weaver, J. Wilde, and L. Zalesky. Euclid Quick Data Release (Q1) – Data release overview, March 2025a. URL <http://arxiv.org/abs/2503.15302>. arXiv:2503.15302 [astro-ph].

Euclid Collaboration, M. Siudek, M. Huertas-Company, M. Smith, G. Martinez-Solaeche, F. Lanusse, S. Ho, E. Angeloudi, P. A. C. Cunha, H. Domínguez Sánchez, M. Dunn, Y. Fu, P. Iglesias-Navarro, J. Junais, J. H. Knapen, B. Laloux, M. Mezcua, W. Roster, G. Stevens, J. Vega-Ferrero, N. Aghanim, B. Altieri, A. Amara, S. Andreon, N. Auricchio, H. Aussel, C. Baccigalupi, M. Baldi, S. Bardelli, P. Battaglia, A. Biviano, A. Bonchi, E. Branchini, M. Brescia, J. Brinchmann, S. Camera, G. Cañas-Herrera, V. Capobianco, C. Carbone, J. Carretero, S. Casas, F. J. Castander, M. Castellano, G. Castignani, S. Cavuoti, K. C. Chambers, A. Cimatti, C. Colodro-Conde, G. Congedo, C. J. Conselice, L. Conversi, Y. Copin, F. Courbin, H. M. Courtois, M. Cropper, A. Da Silva, H. Degaudenzi, G. De Lucia, A. M. Di Giorgio, J. Dinis, C. Dolding, H. Dole, F. Dubath, C. A. J. Duncan, X. Dupac, S. Dusini, S. Escoffier, M. Farina, R. Farinelli, F. Faustini, S. Ferriol, F. Finelli, S. Fotopoulou, M. Frailis, E. Franceschi, S. Galeotta, K. George, B. Gillis, C. Giocoli, J. Gracia-Carpio, B. R. Granett, A. Grazian, F. Grupp, S. Gwyn, S. V. H. Haugan, W. Holmes, I. M. Hook, F. Hormuth, A. Hornstrup, K. Jahnke, M. Jhabvala, E. Keihänen, S. Kermiche, A. Kiessling, B. Kubik, M. Kümmel, M. Kunz, H. Kurki-Suonio, Q. Le Boule, A. M. C. Le Brun, D. Le Mignant, S. Ligori, P. B. Lilje, V. Lindholm, I. Lloro, G. Mainetti, D. Maino, E. Maiorano, O. Mansutti, S. Marcin, O. Marggraf, M. Martinelli, N. Martinet, F. Marulli, R. Massey, S. Maurogordato, H. J. McCracken,

E. Medinaceli, S. Mei, M. Melchior, Y. Mellier, M. Meneghetti, E. Merlin, G. Meylan, A. Mora, M. Moresco, L. Moscardini, R. Nakajima, C. Neissner, S.-M. Niemi, J. W. Nightingale, C. Padilla, S. Paltani, F. Pasian, K. Pedersen, W. J. Percival, V. Pettorino, S. Pires, G. Polenta, M. Poncet, L. A. Popa, L. Pozzetti, F. Raison, A. Renzi, J. Rhodes, G. Riccio, E. Romelli, M. Roncarelli, R. Saglia, Z. Sakr, A. G. Sánchez, D. Sapone, B. Sartoris, J. A. Schewtschenko, P. Schneider, T. Schrabbach, M. Scodeggio, A. Secroun, G. Seidel, M. Seiffert, S. Serrano, P. Simon, C. Sirignano, G. Sirri, L. Stanco, J. Steinwagner, P. Tallada-Crespí, A. N. Taylor, I. Tereno, S. Toft, R. Toledo-Moreo, F. Torradeflot, I. Tutasaus, L. Valenziano, J. Valiviita, T. Vassallo, G. Verdoes Kleijn, A. Veropalumbo, Y. Wang, J. Weller, A. Zacchei, G. Zamorani, F. M. Zerbi, I. A. Zinchenko, E. Zucca, V. Allevato, M. Ballardini, M. Bolzonella, E. Bozzo, C. Burigana, R. Cabanac, A. Cappi, D. Di Ferdinando, J. A. Escartin Vigo, L. Gabarra, J. Martín-Fleitas, S. Matthew, N. Mauri, R. B. Metcalf, A. Pezzotta, M. Pöntinen, C. Porciani, I. Rissò, V. Scottez, M. Sereno, M. Tenti, M. Viel, M. Wiesmann, Y. Akrami, I. T. Andika, S. Anselmi, M. Archidiacono, F. Atrio-Barandela, C. Benoit, K. Benson, D. Bertacca, M. Bethermin, L. Bisigello, A. Blanchard, L. Blot, M. L. Brown, S. Bruton, A. Calabro, B. Camacho Quevedo, F. Caro, C. S. Carvalho, T. Castro, Y. Charles, F. Cogato, A. R. Cooray, O. Cucciati, S. Davini, F. De Paolis, G. Desprez, A. Díaz-Sánchez, J. J. Diaz, S. Di Domizio, J. M. Diego, P.-A. Duc, A. Enia, Y. Fang, A. G. Ferrari, P. G. Ferreira, A. Finoguenov, A. Fontana, A. Franco, K. Ganga, J. García-Bellido, T. Gasparotto, V. Gautard, E. Gaztanaga, F. Giacomini, F. Gianotti, G. Gozaliasl, M. Guidi, C. M. Gutierrez, A. Hall, W. G. Hartley, S. Hemmati, C. Hernández-Monteagudo, H. Hildebrandt, J. Hjorth, J. J. E. Kajava, Y. Kang, V. Kansal, D. Karagiannis, K. Kiiveri, C. C. Kirkpatrick, S. Kruk, J. Le Graet, L. Legrand, M. Lembo, F. Lepori, G. Leroy, G. F. Lesci, J. Lesgourgues, L. Leuzzi, T. I. Liaudat, A. Loureiro, J. Macias-Perez, G. Maggio, M. Magliocchetti, E. A. Magnier, F. Mannucci, R. Maoli, C. J. A. P. Martins, L. Maurin, M. Miluzio, P. Monaco, C. Moretti, G. Morgante, C. Murray, K. Naidoo, A. Navarro-Alsina, S. Nesseris, F. Passalacqua, K. Paterson, L. Patrizii, A. Pisani, D. Potter, S. Quai, M. Radovich, S. Sacquegna, M. Sahlén, D. B. Sanders, E. Sarpa, A. Schneider, D. Sciotti, D. Scognamiglio, E. Sellentin, L. C. Smith, K. Tanidis, G. Testera, R. Teyssier, S. Tosi, A. Troja, M. Tucci, C. Valieri, A. Venhola, D. Vergani, G. Verza, P. Vielzeuf, N. A. Walton, and J. G. Sorce. Euclid Quick Data Release (Q1) Exploring galaxy properties with a multi-modal foundation model, March 2025b. URL <http://arxiv.org/abs/2503.15312>. arXiv:2503.15312 [astro-ph].

Euclid Collaboration, M. Walmsley, M. Huertas-Company, L. Quilley, K. L. Masters, S. Kruk, K. A. Remmelgas, J. J. Popp, E. Romelli, D. O’Ryan, H. J. Dickinson, C. J. Lintott, S. Serjeant, R. J. Smethurst, B. Simmons, J. Shingirai Makechemu, I. L. Garland, H. Roberts, K. Mantha, L. F. Fortson, T. Géron, W. Keel, E. M. Baeten, C. Macmillan, J. Bovy, S. Casas, C. De Leo, H. Domínguez Sánchez, J. Katona, A. Kovács, N. Aghanim, B. Altieri, A. Amara, S. Andreon, N. Auricchio, H. Aussel, C. Baccigalupi, M. Baldi, A. Balestra, S. Bardelli, A. Basset, P. Battaglia, R. Bender, A. Biviano, A. Bonchi, E. Branchini, M. Brescia, J. Brinchmann, S. Camera, G. Cañas-Herrera, V. Capobianco, C. Carbone, J. Carretero, F. J. Castander, M. Castellano, G. Castignani, S. Cavuoti, K. C. Chambers, A. Cimatti, C. Colodro-Conde, G. Congedo, C. J. Conselice, L. Conversi, Y. Copin, F. Courbin, H. M. Courtois, M. Cropper, A. Da Silva, H. Degaudenzi, G. De Lucia, A. M. Di Giorgio, C. Dolding, H. Dole, F. Dubath, C. A. J. Duncan, X. Dupac, S. Dusini, A. Ealet, S. Escoffier, M. Fabricius, M. Farina, R. Farinelli, F. Faustini, F. Finelli, P. Fosalba, S. Fotopoulou, M. Frailis, E. Franceschi, S. Galeotta, K. George, B. Gillis, C. Giocoli, P. Gómez-Alvarez, J. Gracia-Carpio, B. R. Granett, A. Grazian, F. Grupp, S. Gwyn, S. V. H. Haugan, H. Hoekstra, W. Holmes, I. M. Hook, F. Hormuth, A. Hornstrup, P. Hudelot, K. Jahnke, M. Jhabvala, B. Joachimi, E. Keihänen, S. Kermiche, A. Kiessling, R. Kohley, B. Kubik, K. Kuijken, M. Kümmel, M. Kunz, H. Kurki-Suonio, O. Lahav, Q. Le Boule’h, A. M. C. Le Brun, D. Le Mignant, P. Liebing, S. Ligori, P. B. Lilje, V. Lindholm, I. Lloro, G. Mainetti, D. Maino, E. Maiorano, O. Mansutti, S. Marcin, O. Marggraf, M. Martinelli, N. Martinet, F. Marulli, R. Massey, S. Maurogordato, H. J. McCracken, E. Medinaceli, S. Mei, M. Melchior, Y. Mellier, M. Meneghetti, E. Merlin, G. Meylan, A. Mora, M. Moresco, L. Moscardini, R. Nakajima, C. Neissner, R. C. Nichol, S.-M. Niemi, J. W. Nightingale, C. Padilla, S. Paltani, F. Pasian, K. Pedersen, W. J. Percival, V. Pettorino, S. Pires, G. Polenta, M. Poncet, L. A. Popa, L. Pozzetti, F. Raison, R. Rebolo, A. Renzi, J. Rhodes, G. Riccio, M. Roncarelli, B. Rusholme, R. Saglia, Z. Sakr, A. G. Sánchez, D. Sapone, B. Sartoris, J. A. Schewtschenko, P. Schneider, T. Schrabbach, M. Scodeggio, A. Secroun, G. Seidel, M. Seiffert, S. Serrano, P. Simon, C. Sirignano, G. Sirri, L. Stanco, J. Steinwagner, P. Tallada-Crespí, D. Tavagnacco, A. N. Taylor, H. I. Teplitz, I. Tereno, N. Tessore, S. Toft, R. Toledo-Moreo, F. Torradeflot, I. Tutasaus, E. A. Valentijn, L. Valenziano, J. Valiviita, T. Vassallo, G. Verdoes Kleijn, A. Veropalumbo, Y. Wang, J. Weller, A. Zacchei, G. Zamorani, F. M.

Zerbi, I. A. Zinchenko, E. Zucca, V. Allevato, M. Ballardini, M. Bolzonella, E. Bozzo, C. Burigana, R. Cabanac, A. Cappi, D. Di Ferdinando, J. A. Escartin Vigo, L. Gabarra, J. Martín-Fleitas, S. Matthew, N. Mauri, R. B. Metcalf, A. Pezzotta, M. Pöntinen, C. Porciani, I. Risso, V. Scottez, M. Sereno, M. Tenti, M. Viel, M. Wiesmann, Y. Akrami, I. T. Andika, S. Anselmi, M. Archidiacono, F. Atrio-Barandela, C. Benoit, K. Benson, D. Bertacca, M. Bethermin, L. Bisigello, A. Blanchard, L. Blot, H. Böhringer, M. L. Brown, S. Bruton, F. Buitrago, A. Calabro, B. Camacho Quevedo, F. Caro, C. S. Carvalho, T. Castro, F. Cogato, A. R. Cooray, O. Cucciati, S. Davini, F. De Paolis, G. Desprez, A. Díaz-Sánchez, J. J. Diaz, S. Di Domizio, J. M. Diego, P.-A. Duc, A. Enia, Y. Fang, A. G. Ferrari, A. Finoguenov, A. Fontana, A. Franco, K. Ganga, J. García-Bellido, T. Gasparetto, V. Gautard, E. Gaztanaga, F. Giacomini, G. Gozaliasl, M. Guidi, C. M. Gutierrez, A. Hall, W. G. Hartley, S. Hemmati, C. Hernández-Monteagudo, H. Hildebrandt, J. Hjorth, J. J. E. Kajava, Y. Kang, V. Kansal, D. Karagiannis, K. Kiiveri, C. C. Kirkpatrick, J. Le Graet, L. Legrand, M. Lembo, F. Lepori, G. Leroy, G. F. Lesci, J. Lesgourgues, L. Leuzzi, T. I. Liaudat, A. Loureiro, J. Macias-Perez, G. Maggio, M. Magliocchetti, F. Mannucci, R. Maoli, C. J. A. P. Martins, L. Maurin, M. Miluzio, P. Monaco, C. Moretti, G. Morgante, C. Murray, S. Nadathur, K. Naidoo, A. Navarro-Alsina, S. Nesseris, F. Passalacqua, K. Paterson, L. Patrizii, A. Pisani, D. Potter, S. Quai, M. Radovich, P.-F. Rocci, G. Rodighiero, S. Sacquegna, M. Sahlén, D. B. Sanders, E. Sarpa, C. Scarlata, J. Schaye, A. Schneider, M. Schultheis, D. Sciotti, E. Sellentin, F. Shankar, L. C. Smith, K. Tanidis, G. Testera, R. Teyssier, S. Tosi, A. Troja, M. Tucci, C. Valieri, A. Venhola, D. Vergani, G. Verza, P. Vielzeuf, N. A. Walton, E. Soubrie, and D. Scott. Euclid Quick Data Release (Q1): First visual morphology catalogue, March 2025c. URL <http://arxiv.org/abs/2503.15310>. arXiv:2503.15310 [astro-ph].

Amirreza Dolatpour Fathkouhi and Geoffrey Charles Fox. AstroMAE: Redshift Prediction Using a Masked Autoencoder with a Novel Fine-Tuning Architecture, September 2024. URL <http://arxiv.org/abs/2409.01825>. arXiv:2409.01825 [cs].

Christoph Feichtenhofer, Haoqi Fan, Yanghao Li, and Kaiming He. Masked Autoencoders As Spatiotemporal Learners, October 2022. URL <http://arxiv.org/abs/2205.09113>. arXiv:2205.09113 [cs].

Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders, 2024. URL <https://arxiv.org/abs/2406.04093>.

Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut Learning in Deep Neural Networks. *Nature Machine Intelligence*, 2(11):665–673, November 2020. ISSN 2522-5839. doi: 10.1038/s42256-020-00257-z. URL <http://arxiv.org/abs/2004.07780>. arXiv:2004.07780 [cs].

Tom George Grigg, Dan Busbridge, Jason Ramapuram, and Russ Webb. Do self-supervised and supervised methods learn similar visual representations?, 2021. URL <https://arxiv.org/abs/2110.00528>.

Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked Autoencoders Are Scalable Vision Learners. In 2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 15979–15988, June 2022. doi: 10.1109/CVPR52688.2022.01553. URL <https://ieeexplore.ieee.org/document/9879206>. ISSN: 2575-7075.

Adam Karvonen, Can Rager, Johnny Lin, Curt Tigges, Joseph Bloom, David Chanin, Yeu-Tong Lau, Eoin Farrell, Callum McDougall, Kola Ayonrinde, Demian Till, Matthew Wearden, Arthur Conmy, Samuel Marks, and Neel Nanda. SAE Bench: A Comprehensive Benchmark for Sparse Autoencoders in Language Model Interpretability, June 2025. URL <http://arxiv.org/abs/2503.09532>. arXiv:2503.09532 [cs].

M. Lochner and B. A. Bassett. ASTRONOMALY: Personalised active anomaly detection in astronomical data. *Astronomy and Computing*, 36:100481, July 2021. doi: 10.1016/j.ascom.2021.100481.

Alireza Makhzani and Brendan Frey. k-sparse autoencoders, 2014. URL <https://arxiv.org/abs/1312.5663>.

Samuel Marks, Johannes Treutlein, Trenton Bricken, Jack Lindsey, Jonathan Marcus, Siddharth Mishra-Sharma, Daniel Ziegler, Emmanuel Ameisen, Joshua Batson, Tim Belonax, Samuel R. Bowman, Shan Carter, Brian Chen, Hoagy Cunningham, Carson Denison, Florian Dietz, Satvik Golechha, Akbir Khan, Jan Kirchner, Jan Leike, Austin Meek, Kei Nishimura-Gasparian, Euan Ong, Christopher Olah, Adam Pearce, Fabien Roger, Jeanne Salle, Andy Shih, Meg Tong, Drake

- Thomas, Kelley Rivoire, Adam Jermyn, Monte MacDiarmid, Tom Henighan, and Evan Hubinger. Auditing language models for hidden objectives, 2025. URL <https://arxiv.org/abs/2503.10965>.
- Rajiv Movva, Kenny Peng, Nikhil Garg, Jon Kleinberg, and Emma Pierson. Sparse autoencoders for hypothesis generation, 2025. URL <https://arxiv.org/abs/2502.04382>.
- Aashiq Muhamed, Mona Diab, and Virginia Smith. Decoding dark matter: Specialized sparse autoencoders for interpreting rare concepts in foundation models, 2024. URL <https://arxiv.org/abs/2411.00743>.
- Charles O'Neill, Christine Ye, Kartheik Iyer, and John F. Wu. Disentangling Dense Embeddings with Sparse Autoencoders. *arXiv e-prints*, art. arXiv:2408.00657, August 2024. doi: 10.48550/arXiv.2408.00657.
- David O'Ryan and Pablo Gómez. Identifying astrophysical anomalies in 99.6 million cutouts from the hubble legacy archive using anomalousmatch, 2025. URL <https://arxiv.org/abs/2505.03508>.
- Kenny Peng, Rajiv Movva, Jon Kleinberg, Emma Pierson, and Nikhil Garg. Use sparse autoencoders to discover unknown concepts, not to act on known concepts, 2025. URL <https://arxiv.org/abs/2506.23845>.
- Colorado J. Reed, Ritwik Gupta, Shufan Li, Sarah Brockman, Christopher Funk, Brian Clipp, Kurt Keutzer, Salvatore Candido, Matt Uyttendaele, and Trevor Darrell. Scale-MAE: A Scale-Aware Masked Autoencoder for Multiscale Geospatial Representation Learning, September 2023. URL <http://arxiv.org/abs/2212.14532>. arXiv:2212.14532 [cs].
- R. Scaramella, J. Amiaux, Y. Mellier, C. Burigana, C. S. Carvalho, J.-C. Cuillandre, A. Da Silva, A. Derosa, J. Dinis, E. Maiorano, M. Maris, I. Tereno, R. Laureijs, T. Boenke, G. Buenadicha, X. Dupac, L. M. Gaspar Venancio, P. Gómez-Álvarez, J. Hoar, J. Lorenzo Alvarez, G. D. Racca, G. Saavedra-Criado, J. Schwartz, R. Vavrek, M. Schirmer, H. Aussel, R. Azzollini, V. F. Cardone, M. Cropper, A. Ealet, B. Garilli, W. Gillard, B. R. Granett, L. Guzzo, H. Hoekstra, K. Jahnke, T. Kitching, M. Meneghetti, L. Miller, R. Nakajima, S. M. Niemi, F. Pasian, W. J. Percival, M. Sauvage, M. Scodellaggio, S. Wachter, A. Zacchei, N. Aghanim, A. Amara, T. Auphan, N. Auricchio, S. Awan, A. Balestra, R. Bender, C. Bodendorf, D. Bonino, E. Branchini, S. Brau-Nogue, M. Brescia, G. P. Candini, V. Capobianco, C. Carbone, R. G. Carlberg, J. Carretero, R. Casas, F. J. Castander, M. Castellano, S. Cavuoti, A. Cimatti, R. Cledassou, G. Congedo, C. J. Conselice, L. Conversi, Y. Copin, L. Corcione, A. Costille, F. Courbin, H. Degaudenzi, M. Douspis, F. Dubath, C. A. J. Duncan, S. Dusini, S. Farrens, S. Ferriol, P. Fosalba, N. Fourmanoit, M. Frailis, E. Franceschi, P. Franzetti, M. Fumana, B. Gillis, C. Giocoli, A. Grazian, F. Grupp, S. V. H. Haugan, W. Holmes, F. Hormuth, P. Hudelot, S. Kermiche, A. Kiessling, M. Kilbinger, R. Kohley, B. Kubik, M. Kümmel, M. Kunz, H. Kurki-Suonio, S. Ligori, P. B. Lilje, I. Lloro, O. Mansutti, O. Marggraf, K. Markovic, F. Marulli, R. Massey, S. Maurogordato, M. Melchior, E. Merlin, G. Meylan, J. J. Mohr, M. Moresco, B. Morin, L. Moscardini, E. Munari, R. C. Nichol, C. Padilla, S. Paltani, J. Peacock, K. Pedersen, V. Pettorino, S. Pires, M. Ponchet, L. Popa, L. Pozzetti, F. Raison, R. Rebolo, J. Rhodes, H.-W. Rix, M. Roncarelli, E. Rossetti, R. Saglia, P. Schneider, T. Schrabback, A. Secroun, G. Seidel, S. Serrano, C. Sirignano, G. Sirri, J. Skottfelt, L. Stanco, J. L. Starck, P. Tallada-Crespí, D. Tavagnacco, A. N. Taylor, H. I. Teplitz, R. Toledo-Moreo, F. Torradeflot, M. Trifoglio, E. A. Valentijn, L. Valenziano, G. A. Verdoes Kleijn, Y. Wang, N. Welikala, J. Weller, M. Wetzstein, G. Zamorani, J. Zoubian, S. Andreon, M. Baldi, S. Bardelli, A. Boucaud, S. Camera, G. Fabbian, R. Farinelli, J. Graciá-Carpio, D. Maino, E. Medinaceli, S. Mei, C. Neissner, G. Polenta, A. Renzi, E. Romelli, C. Rosset, F. Sureau, M. Tenti, T. Vassallo, E. Zucca, C. Baccigalupi, A. Balaguera-Antolínez, P. Battaglia, A. Biviano, S. Borgani, E. Bozzo, R. Cabanac, A. Cappi, S. Casas, G. Castignani, C. Colodro-Conde, J. Coupon, H. M. Courtois, J. Cuby, S. de la Torre, S. Desai, D. Di Ferdinando, H. Dole, M. Fabricius, M. Farina, P. G. Ferreira, F. Finelli, P. Flose-Reimberg, S. Fotopoulou, S. Galeotta, K. Ganga, G. Gozaliasl, I. M. Hook, E. Keihanen, C. C. Kirkpatrick, P. Liebing, V. Lindholm, G. Mainetti, M. Martinelli, N. Martinet, M. Maturi, H. J. McCracken, R. B. Metcalf, G. Morgante, J. Nightingale, A. Nucita, L. Patrizii, D. Potter, G. Riccio, A. G. Sánchez, D. Sapone, J. A. Schewtschenko, M. Schultheis, V. Scottez, R. Teyssier, I. Tutasaus, J. Valiviita, M. Viel, W. Vriend, and L. Whittaker. Euclid preparation: I. The Euclid Wide Survey, August 2021. URL <http://arxiv.org/abs/2108.01201>. arXiv:2108.01201.
- D. Spergel, N. Gehrels, C. Baltay, D. Bennett, J. Breckinridge, M. Donahue, A. Dressler, B. S. Gaudi, T. Greene, O. Guyon, C. Hirata, J. Kalirai, N. J. Kasdin, B. Macintosh, W. Moos, S. Perlmutter,

M. Postman, B. Rauscher, J. Rhodes, Y. Wang, D. Weinberg, D. Benford, M. Hudson, W. S. Jeong, Y. Mellier, W. Traub, T. Yamada, P. Capak, J. Colbert, D. Masters, M. Penny, D. Savransky, D. Stern, N. Zimmerman, R. Barry, L. Bartusek, K. Carpenter, E. Cheng, D. Content, F. Dekens, R. Demers, K. Grady, C. Jackson, G. Kuan, J. Kruk, M. Melton, B. Nemati, B. Parvin, I. Poberezhskiy, C. Peddie, J. Ruffa, J. K. Wallace, A. Whipple, E. Wollack, and F. Zhao. Wide-field infrared survey telescope-astrophysics focused telescope assets wfirst-afta 2015 report, 2015. URL <https://arxiv.org/abs/1503.03757>.

John F. Wu. Insights into Galaxy Evolution from Interpretable Sparse Feature Networks. *The Astrophysical Journal*, 980(2):183, February 2025. doi: 10.3847/1538-4357/adadec. URL <https://iopscience.iop.org/article/10.3847/1538-4357/adadec>.

Christine Ye, Charles O'Neill, John F Wu, and Kartheik G. Iyer. Steering semantic search with interpretable features from sparse autoencoders. In *MINT: Foundation Model Interventions*, 2024. URL <https://openreview.net/forum?id=oacksuh7Tu>.

A Interpreting SAE features beyond the top 64

In Figure 4, we show that SAE features beyond the top 64 are still aligned with GZ. For the supervised case (left-hand panel), the correlation grows weaker as we reach the later groups in the Matryoshka SAE: typical correlations for the top 64 features are ~ 0.3 , whereas mean Spearman correlations in the final group with GZ are near ~ 0.1 .

Meanwhile, the SAE features for self-supervised embeddings (right-hand panel) all have a moderately high alignment with GZ features. Intriguingly, this does not vary much with the Matryoshka groups, suggesting that the basis set of MAE features is independent of GZ.

We also show PCA again, which exhibits significantly lower alignment with GZ than the SAE features in both the supervised and self-supervised cases first principal component, which has maximum $|r| \approx 0.87$ and separates smooth galaxies from disk or featured galaxies). Principal components beyond the top 64 have nearly zero max Spearman correlation r with GZ features (and are consistent with random), so we omit them from Figure 4.

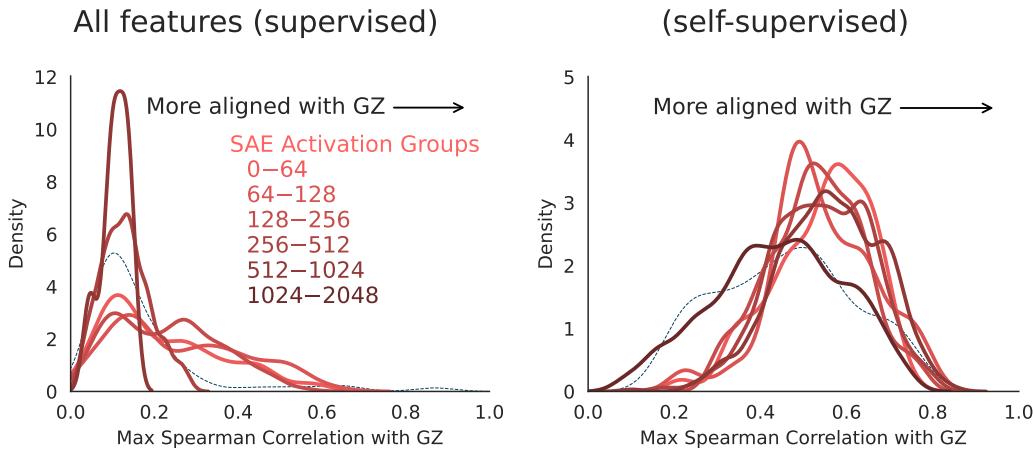


Figure 4: How aligned are features with GZ, just like in Figure 3? We show results for both supervised (*left*) and self-supervised (*right*) approaches. This time we use all of the SAE activations, binned by their Matryoshka group sizes (progressing from lighter to darker red). In all panels, the top 64 elements from PCA are shown in a thin dashed blue line.

B Feature Predictability

In the same spirit as the “Feature Alignment” analysis, we can also fit linear regression to predict each learned feature from all Galaxy Zoo labels and measure R^2 . This is not a trivial deprojection for the supervised case because the GZ model has a non-linear classification head.

The R^2 metric probes how *predictable* any SAE feature is, based on known GZ features. If we are unable to explain the feature variance using existing morphology classifications, then it may suggest that the feature is novel. (It could also signify that the feature is highly polysemantic!)

In Figure 5, we show how well GZ classes can predict the strengths of SAE activations and principal components. We find qualitatively similar results to before: for the supervised case, the PCA features are generally difficult to predict (aside from the first handful), and the SAE activations also become harder to predict as we progress to later Matryoshka groups. For the MAE embeddings, the SAE features are moderately predictable (and PCA features less so). Again, we suspect this is because MAE embeddings contain salient information about image artifacts, and GZ vote fractions can partially reconstruct these features.

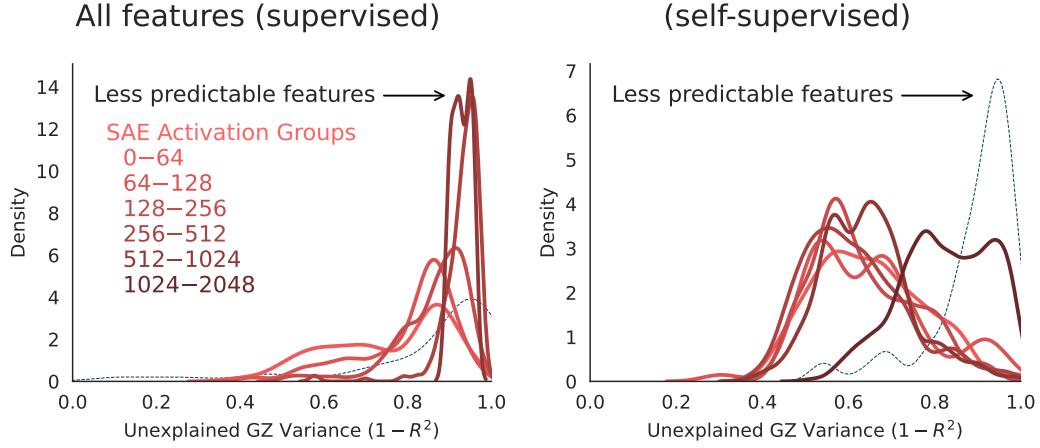


Figure 5: How unpredictable—or perhaps, “novel”—are our features via linear combinations of GZ vote fractions? Figure in same style as Figure 4.

C Expected Spearman Correlations with Random Features

We can fashion a null test by sampling random features and computing the expectation value of the maximum Spearman r by using extreme value statistics: $E[\max(r)] \approx \sqrt{2 \log(2c)/(n-1)}$, where we have $c = 46$ GZ classes, $k = 64$ random features, and $n = 2378$ or 76061 test set samples respectively in the supervised and self-supervised test sets. This approximation is valid in the limit of moderate to large n , and we use $2c$ rather than c because we take the absolute Spearman correlation (probing a two-tailed distribution). The standard deviation of the maximum Spearman rank can be approximated via $\pi/\sqrt{k(n-1)\log(2c)}$. These expectation values approximate to 0.041 ± 0.002 for the supervised case, and 0.0072 ± 0.0003 for self-supervised model embeddings.